

ພວດີເປີ່ງເຫັນວ່າ pdf ທີ່ສົ່ງໄປກ່ອນໜ້າ (ສົ່ງທັນ) ມັນບໍ່ຄໍາໃຫ້ຄໍາຕອນນາງຂ້ອຍໄປເລຍຄົນ ເລຍໝາດ້ວເຊີ້ນສົ່ງມາອີກຮອນ ຄ້າໄມ້ທັກຄະແນນສົ່ງສ່າຍຈະດີໃຈມາກຄົນ 

Precipitation Nowcasting using Neural Networks

In this exercise, you are going to build a set of deep learning models on a real world task using PyTorch. PyTorch is an open source machine learning framework based on the Torch library, used for applications such as computer vision and natural language processing, primarily developed by Facebook's AI Research lab (FAIR).

Setting up to use the gpu

Before we start, we need to change the environment of Colab to use GPU. Do so by:

Runtime -> Change runtime type -> Hardware accelerator -> GPU

Deep Neural Networks with PyTorch

To complete this exercise, you will need to build deep learning models for precipitation nowcasting. You will build a subset of the models shown below:

- Fully Connected (Feedforward) Neural Network
- Two-Dimentional Convolution Neural Network (2D-CNN)
- Recurrent Neural Network with Gated Recurrent Unit (GRU)

and one more model of your choice to achieve the highest score possible.

We provide the code for data cleaning and some starter code for PyTorch in this notebook but feel free to modify those parts to suit your needs. Feel free to use additional libraries (e.g. scikit-learn) as long as you have a model for each type mentioned above.

This notebook assumes you have already installed PyTorch with python3 and had GPU enabled. If you run this exercise on Colab you are all set.

Precipitation Nowcasting

Precipitation nowcasting is the task of predicting the amount of rainfall in a certain region given some kind of sensor data. The term nowcasting refers to tasks that try to predict the current or near future conditions (within 6 hours).

You will be given satellite images in 3 different bands covering a 5 by 5 region from different parts of Thailand. In other words, your input will be a 5x5x3 image. Your task is to predict the amount of rainfall in the center pixel. You will first do the prediction using just a simple fully-connected neural network that view each pixel as different input features.

Since your input is basically an image, we will then view the input as an image and apply CNN to do the prediction. Finally, we can also add a time component since weather prediction can benefit greatly using previous time frames. Each data point actually contain 5 time steps, so each input data point has a size of 5x5x5x3 (time x height x width x channel), and the output data has a size of 5 (time). You will use this time information when you work with RNNs.

Finally, we would like to thank the Thai Meteorological Department for providing the data for this assignment.

```
In [1]: # !nvidia-smi
```

```
In [2]: # # For summarizing and visualizing models  
# !pip install torchinfo  
# !pip install torchviz
```

Weights and Biases

[Weights and Biases \(https://docs.wandb.ai/company\)](https://docs.wandb.ai/company) (wandb) is an experiment tracking tool for machine learning. It can log and visualize experiments in real time. It supports many popular ML frameworks, and obviously PyTorch is one of them. In this notebook you will learn how to log general metrics like losses, parameter distributions, and gradient distribution with wandb.

To install wandb, run the cell below

```
In [3]: # !pip install wandb
```

Setup

1. Register [Wandb account \(https://wandb.ai/login?signup=true\)](https://wandb.ai/login?signup=true) (and confirm your email)
2. wandb login and copy paste the API key when prompt

```
In [4]: # !wandb login
```

```
In [1]: import os
import numpy as np
import pickle
import pandas as pd
import matplotlib.pyplot as plt
import urllib
import wandb
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.transforms as transforms

from sklearn import preprocessing
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from torchinfo import summary
from tqdm.notebook import tqdm

torch.__version__ # 1.10.0+cu111
```

Out[1]: '2.2.1'

```
In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

Out[2]: device(type='cuda')

Loading the data

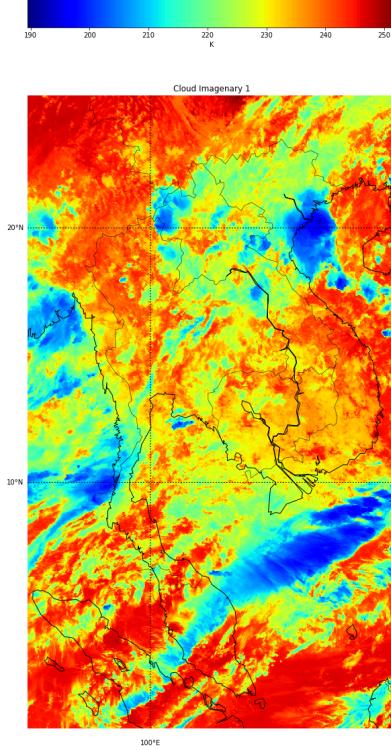
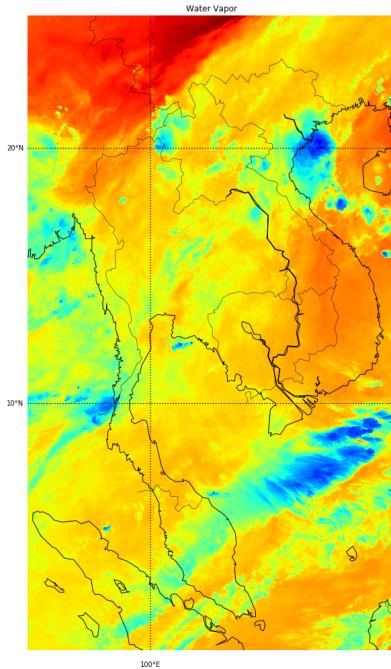
Get the data set by going [here \(<https://drive.google.com/file/d/1NWR22fVVE0tO2Q5EbaPPrRKPhUem-jbw/view?usp=sharing>\)](https://drive.google.com/file/d/1NWR22fVVE0tO2Q5EbaPPrRKPhUem-jbw/view?usp=sharing) and click add to drive.

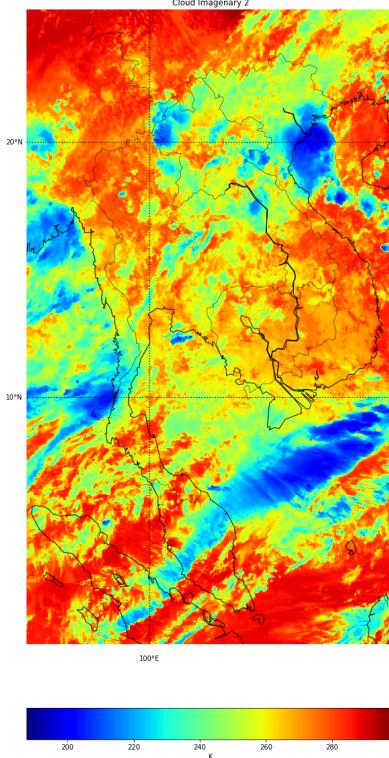
```
In [6]: # from google.colab import drive
# drive.mount('/content/gdrive/')
```

```
In [7]: # !tar -xvf '/content/gdrive/My Drive/nowcastingHWdataset.tar.gz'
```

Data Explanation

The data is an hourly measurement of water vapor in the atmosphere, and two infrared measurements of cloud imagery on a latitude-longitude coordinate. Each measurement is illustrated below as an image. These three features are included as different channels in your input data.





We also provide the hourly precipitation (rainfall) records in the month of June, July, August, September, and October from weather stations spreaded around the country. A 5x5 grid around each weather station at a particular time will be paired with the precipitation recorded at the corresponding station as input and output data. Finally, five adjacent timesteps are stacked into one sequence.

The month of June-August are provided as training data, while the months of September and October are used as validation and test sets, respectively.

Reading data

```
In [3]: def read_data(months, data_dir='dataset'):
    features = np.array([], dtype=np.float32).reshape(0,5,5,5,3)
    labels = np.array([], dtype=np.float32).reshape(0,5)
    for m in months:
        filename = 'features-{}.pk'.format(m)
        with open(os.path.join(data_dir,filename), 'rb') as file:
            features_temp = pickle.load(file)
        features = np.concatenate((features, features_temp), axis=0)

        filename = 'labels-{}.pk'.format(m)
        with open(os.path.join(data_dir,filename), 'rb') as file:
            labels_temp = pickle.load(file)
        labels = np.concatenate((labels, labels_temp), axis=0)

    return features, labels
```

```
In [4]: # use data from month 6,7,8 as training set
x_train, y_train = read_data(months=[6,7,8])

# use data from month 9 as validation set
x_val, y_val = read_data(months=[9])

# use data from month 10 as test set
x_test, y_test = read_data(months=[10])

print('x_train shape:',x_train.shape)
print('y_train shape:', y_train.shape, '\n')
print('x_val shape:',x_val.shape)
print('y_val shape:', y_val.shape, '\n')
print('x_test shape:',x_test.shape)
print('y_test shape:', y_test.shape)
```

```
x_train shape: (229548, 5, 5, 5, 3)
y_train shape: (229548, 5)
```

```
x_val shape: (92839, 5, 5, 5, 3)
y_val shape: (92839, 5)
```

```
x_test shape: (111715, 5, 5, 5, 3)
y_test shape: (111715, 5)
```

features

- dim 0: number of entries
- dim 1: number of time-steps in ascending order
- dim 2,3: a 5x5 grid around rain-measured station
- dim 4: water vapor and two cloud imangenaries

labels

- dim 0: number of entries
- dim 1: number of precipitation for each time-step

Three-Layer Feedforward Neural Networks

```
In [5]: # Dataset need to be reshaped to make it suitable for feedforward model
def preprocess_for_ff(x_train, y_train, x_val, y_val):
    x_train_ff = x_train.reshape((-1, 5*5*3))
    y_train_ff = y_train.reshape((-1, 1))
    x_val_ff = x_val.reshape((-1, 5*5*3))
    y_val_ff = y_val.reshape((-1, 1))
    x_test_ff = x_test.reshape((-1, 5*5*3))
    y_test_ff = y_test.reshape((-1, 1))

    return x_train_ff, y_train_ff, x_val_ff, y_val_ff, x_test_ff, y_test_ff

x_train_ff, y_train_ff, x_val_ff, y_val_ff, x_test_ff, y_test_ff = preprocess_
for_ff(x_train, y_train, x_val, y_val)
print(x_train_ff.shape, y_train_ff.shape)
print(x_val_ff.shape, y_val_ff.shape)
print(x_test_ff.shape, y_test_ff.shape)

(1147740, 75) (1147740, 1)
(464195, 75) (464195, 1)
(558575, 75) (558575, 1)
```

TODO#1

Explain each line of code in the function preprocess_for_ff()

Ans: flatten data.

Dataset

To prepare a DataLoader in order to feed data into the model, we need to create a `torch.utils.data.Dataset` object first. (Learn more about it [here](https://pytorch.org/docs/stable/data.html#map-style-datasets) (<https://pytorch.org/docs/stable/data.html#map-style-datasets>))

Dataset is a simple class that the DataLoader will get data from, most of its functionality comes from `__getitem__(self, index)` method, which will return a single data point (both input and label). In real world scenarios the method can do some other stuffs such as

1. Load images

If your input (x) are images. Oftentimes you won't be able to fit all the training images into your RAM. Thus, you should pass an array (or list) of image path into the dataloader, and the `__getitem__` will be the one who dynamically loads the actual image from the harddisk for you.

1. Data Normalization

Data normalization helps improve stability of training. Unnormalized data can cause gradients to explode. There are many variants of normalization, but in this notebook we will use either minmax or z-score (std) normalization. Read [this](https://developers.google.com/machine-learning/data-prep/transform/normalization) (<https://developers.google.com/machine-learning/data-prep/transform/normalization>) (or google) if you wish to learn more about data normalization.

1. Data Augmentation

In computer vision, you might want to apply small changes to the images you use in training (adjust brightness, contrast, rotation) so that the model will generalize better on unseen data. There are two kinds of augmentation: static and dynamic. Static augmentation will augment images and save to disk as a new dataset. On the other hand, rather than applying the change initially and use the same change on each image every epoch, dynamic augmentation will augment each data differently for each epoch. Note that augmentation is usually done on the CPU and you might be bounded by the CPU instead. PyTorch has a dedicated [documentation about data augmentation](https://pytorch.org/vision/master/transforms.html) (<https://pytorch.org/vision/master/transforms.html>) if you want to know more.

```
In [6]: class RainfallDatasetFF(Dataset):
    def __init__(self, x, y, normalizer):
        self.x = x.astype(np.float32)
        self.y = y.astype(np.float32)
        self.normalizer = normalizer
        print(self.x.shape)
        print(self.y.shape)

    def __getitem__(self, index):
        x = self.x[index] # Retrieve data
        x = self.normalizer.transform(x.reshape(1, -1)) # Normalize
        y = self.y[index]
        return x, y

    def __len__(self):
        return self.x.shape[0]
```

```
In [7]: def normalizer_std(X):
    scaler = preprocessing.StandardScaler().fit(X)
    return scaler
```

```
def normalizer_minmax(X):
    scaler = preprocessing.MinMaxScaler().fit(X)
    return scaler
```

```
In [8]: normalizer = normalizer_std(x_train_ff) # We will normalize everything based on x_train
```

```
train_dataset = RainfallDatasetFF(x_train_ff, y_train_ff, normalizer)
val_dataset = RainfallDatasetFF(x_val_ff, y_val_ff, normalizer)
test_dataset = RainfallDatasetFF(x_test_ff, y_test_ff, normalizer)
```

```
(1147740, 75)
(1147740, 1)
(464195, 75)
(464195, 1)
(558575, 75)
(558575, 1)
```

DataLoader

DataLoader feeds data from our dataset into the model. We can freely customize batch size, data shuffle for each data split, and much more with DataLoader class. If you're curious about what can you do with PyTorch's DataLoader, you can check [this documentation](https://pytorch.org/docs/stable/data.html) (<https://pytorch.org/docs/stable/data.html>)

```
In [9]: train_loader = DataLoader(train_dataset, batch_size=1024, shuffle=True, pin_memory=True)
val_loader = DataLoader(val_dataset, batch_size=1024, shuffle=False, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=1024, shuffle=False, pin_memory=True)
```

Loss Function

PyTorch has many loss functions readily available for use. We can also write our own custom loss function as well. But for now, we will use [PyTorch's built-in mean squared error loss](https://pytorch.org/docs/stable/generated/torch.nn.MSELoss.html) (<https://pytorch.org/docs/stable/generated/torch.nn.MSELoss.html>).

```
In [10]: loss_fn = nn.MSELoss()
```

TODO#2

Why is the loss MSE?

Ans: Because, this is regression task, and it is easy to find derivative.

Device

Unlike Tensorflow/Keras, PyTorch allows user to freely put any Tensor or objects (loss functions, models, optimizers, etc.) in CPU or GPU. By default, all objects created will be in CPU. In order to use GPU we will have to supply `device = torch.device("cuda")` into the objects to move it to GPU. You will usually see the syntax like `object.to(device)` for moving CPU object to GPU, or `o = Object(..., device=device)` to create the object in the GPU.

```
In [11]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
Out[11]: device(type='cuda')
```

Model

Below, the code for creating a 3-layers fully connected neural network in PyTorch is provided. Run the code and make sure you understand what you are doing. Then, report the results.

```
In [60]: class FeedForwardNN(nn.Module):
    def __init__(self, hidden_size=200):
        super(FeedForwardNN, self).__init__()
        self.ff1 = nn.Linear(75, hidden_size)
        self.ff2 = nn.Linear(hidden_size, hidden_size)
        self.ff3 = nn.Linear(hidden_size, hidden_size)
        self.out = nn.Linear(hidden_size, 1)

    def forward(self, x):
        hd1 = F.relu(self.ff1(x))
        hd2 = F.relu(self.ff2(hd1))
        y = F.relu(self.ff3(hd2))
        y = self.out(y)
        return y.reshape(-1, 1)
```

TODO#3

What is the activation function in the final dense layer? and why? Do you think there is a better activation function for the final layer?

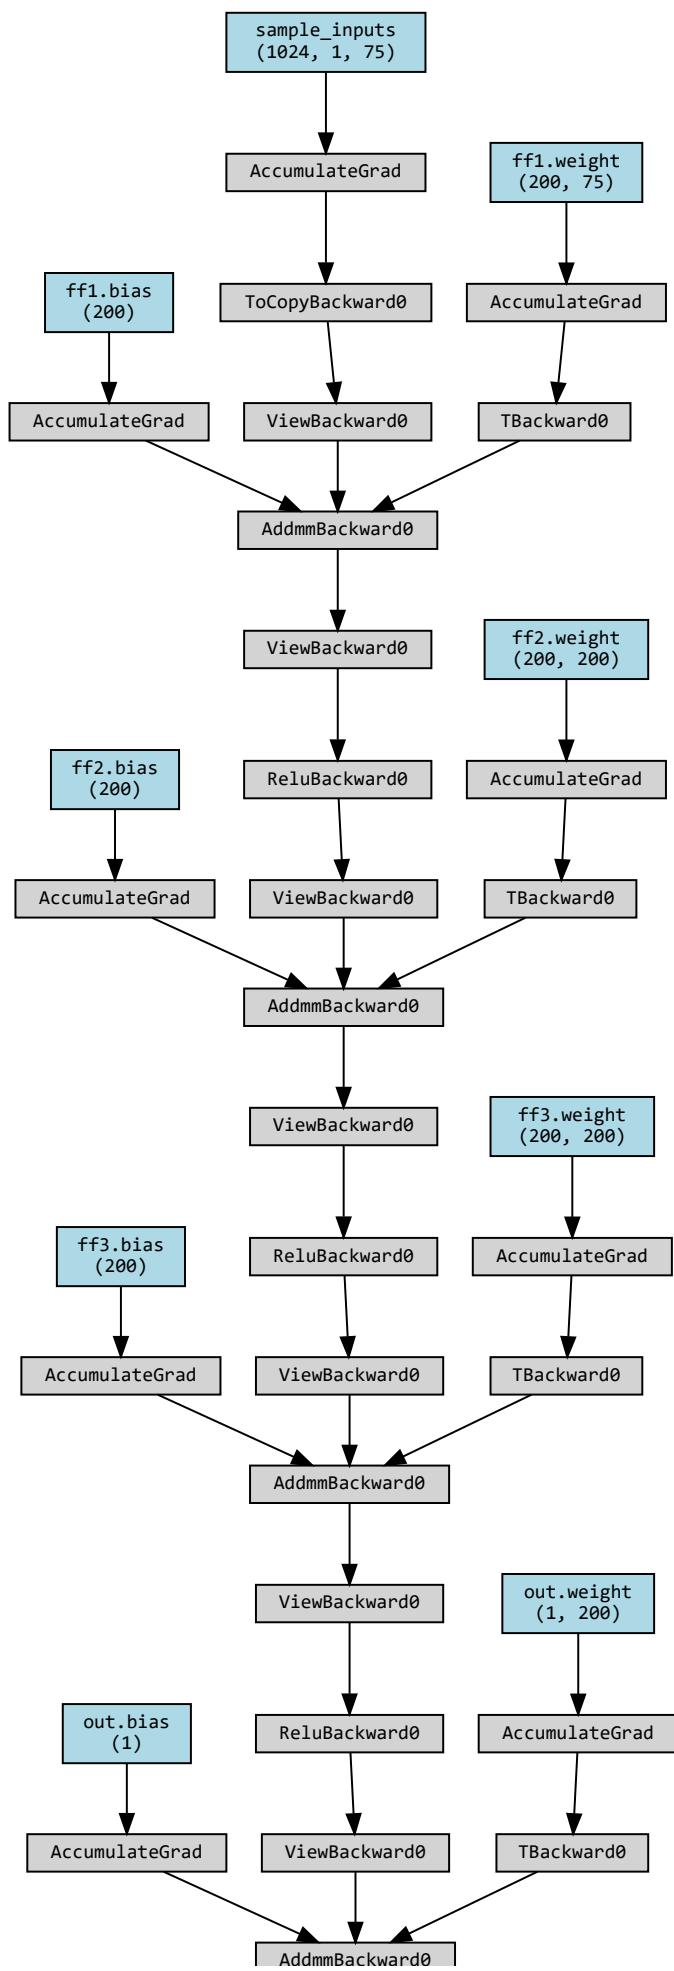
Ans: ReLU, because atmospheric water vapor should not less than 0. So, I think ReLU is the best.

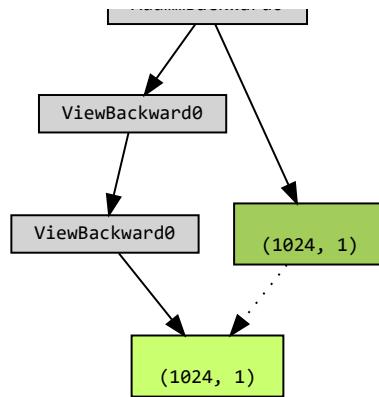
```
In [61]: # Hyperparameters and other configs
config = {
    'architecture': 'feedforward',
    'lr': 0.01,
    'hidden_size': 200,
    'scheduler_factor': 0.2,
    'scheduler_patience': 2,
    'scheduler_min_lr': 1e-4,
    'epochs': 10
}

# Model
model_ff = FeedForwardNN(hidden_size=config['hidden_size'])
model_ff = model_ff.to(device)
optimizer = torch.optim.Adam(model_ff.parameters(), lr=config['lr'])
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    'min',
    factor=config['scheduler_factor'],
    patience=config['scheduler_patience'],
    min_lr=config['scheduler_min_lr']
)
```

```
In [14]: from torchviz import make_dot
# Visualize model with torchviz
sample_inputs = next(iter(train_loader))[0].requires_grad_(True)
sample_y = model_ff(sample_inputs.to(device))
make_dot(sample_y, params=dict(list(model_ff.named_parameters())+(['sample_inputs', sample_inputs])))
```

Out[14]:





In [15]: `summary(model_ff, input_size=(1024, 75))`

Out[15]:

```

=====
=====
Layer (type:depth-idx)          Output Shape      Param #
=====
=====
FeedForwardNN                   [1024, 1]        --
|---Linear: 1-1                 [1024, 200]      15,200
|---Linear: 1-2                 [1024, 200]      40,200
|---Linear: 1-3                 [1024, 200]      40,200
|---Linear: 1-4                 [1024, 1]       201
=====
=====
Total params: 95,801
Trainable params: 95,801
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 98.10
=====
=====
Input size (MB): 0.31
Forward/backward pass size (MB): 4.92
Params size (MB): 0.38
Estimated Total Size (MB): 5.61
=====
=====
```

TODO#4

Explain why the first linear layer has number of parameters = 15200

Ans: Because, input has 75 features map to 200 node in first layer. Therefore, it has $75 \times 200 = 15200$ parameters.

Training

```
In [19]: train_losses = []
val_losses = []
learning_rates = []

# Start wandb run
wandb.init(
    project='precipitation-nowcasting',
    config=config,
)

# Log parameters and gradients
wandb.watch(model_ff, log='all')

for epoch in range(config['epochs']): # Loop over the dataset multiple times

    # Training
    train_loss = []
    current_lr = optimizer.param_groups[0]['lr']
    learning_rates.append(current_lr)

    # Flag model as training. Some layers behave differently in training and
    # inference modes, such as dropout, BN, etc.
    model_ff.train()

    print(f"Training epoch {epoch+1}...")
    print(f"Current LR: {current_lr}")

    for i, (inputs, y_true) in enumerate(tqdm(train_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Reset the gradient
        optimizer.zero_grad()

        # Predict
        y_pred = model_ff(inputs)

        # Calculate loss
        loss = loss_fn(y_pred, y_true)

        # Compute gradient
        loss.backward()

        # Update parameters
        optimizer.step()

        # Log stuff
        train_loss.append(loss)

    avg_train_loss = torch.stack(train_loss).mean().item()
    train_losses.append(avg_train_loss)

    print(f"Epoch {epoch+1} train loss: {avg_train_loss:.4f}")

# Validation
```

```

model_ff.eval()
with torch.no_grad(): # No gradient is required during validation
    print(f"Validating epoch {epoch+1}")
    val_loss = []
    for i, (inputs, y_true) in enumerate(tqdm(val_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Predict
        y_pred = model_ff(inputs)

        # Calculate loss
        loss = loss_fn(y_pred, y_true)

        # Log stuff
        val_loss.append(loss)

    avg_val_loss = torch.stack(val_loss).mean().item()
    val_losses.append(avg_val_loss)
    print(f"Epoch {epoch+1} val loss: {avg_val_loss:.4f}")

    # LR adjustment with scheduler
    scheduler.step(avg_val_loss)

    # Save checkpoint if val_loss is the best we got
    best_val_loss = np.inf if epoch == 0 else min(val_losses[:-1])
    if avg_val_loss < best_val_loss:
        # Save whatever you want
        state = {
            'epoch': epoch,
            'model': model_ff.state_dict(),
            'optimizer': optimizer.state_dict(),
            'scheduler': scheduler.state_dict(),
            'train_loss': avg_train_loss,
            'val_loss': avg_val_loss,
            'best_val_loss': best_val_loss,
        }

        print(f"Saving new best model..")
        torch.save(state, 'model_ff2.pth.tar')

wandb.log({
    'train_loss': avg_train_loss,
    'val_loss': avg_val_loss,
    'lr': current_lr,
})

wandb.finish()
print('Finished Training')

```

Failed to detect the name of this notebook, you can set it manually with the `WANDB_NOTEBOOK_NAME` environment variable to enable code saving.
`wandb`: Currently logged in as: `demonstem`. Use ``wandb login --relogin`` to force relogin

Tracking run with wandb version 0.16.4

Run data is saved locally in `c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240319_235818-4q5utifx`

Syncing run [chocolate-sunset-10 \(https://wandb.ai/demonstem/precipitation-nowcasting/runs/4q5utifx\)](https://wandb.ai/demonstem/precipitation-nowcasting/runs/4q5utifx) to [Weights & Biases \(https://wandb.ai/demonstem/precipitation-nowcasting\)](https://wandb.ai/demonstem/precipitation-nowcasting) ([docs \(https://wandb.me/run\)](https://wandb.me/run))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting> (<https://wandb.ai/demonstem/precipitation-nowcasting>).

View run at <https://wandb.ai/demonstem/precipitation-nowcasting/runs/4q5utifx> (<https://wandb.ai/demonstem/precipitation-nowcasting/runs/4q5utifx>).

Training epoch 1...
Current LR: 0.01

Epoch 1 train loss: 1.9268
Validating epoch 1

Epoch 1 val loss: 1.6583
Saving new best model..
Training epoch 2...
Current LR: 0.01

Epoch 2 train loss: 1.9210
Validating epoch 2

Epoch 2 val loss: 1.6590
Training epoch 3...
Current LR: 0.01

Epoch 3 train loss: 1.9232
Validating epoch 3

Epoch 3 val loss: 1.6614
Training epoch 4...
Current LR: 0.01

Epoch 4 train loss: 1.9234
Validating epoch 4

Epoch 4 val loss: 1.6632
Training epoch 5...
Current LR: 0.002

Epoch 5 train loss: 1.9229
Validating epoch 5

Epoch 5 val loss: 1.6587
Training epoch 6...
Current LR: 0.002

Epoch 6 train loss: 1.9195
Validating epoch 6

Epoch 6 val loss: 1.6577
Saving new best model..
Training epoch 7...
Current LR: 0.002

Epoch 7 train loss: 1.9185
Validating epoch 7

Epoch 7 val loss: 1.6566
Saving new best model..
Training epoch 8...
Current LR: 0.002

Epoch 8 train loss: 1.9182
Validating epoch 8

Epoch 8 val loss: 1.6566
Saving new best model..
Training epoch 9...
Current LR: 0.002

Epoch 9 train loss: 1.9181
Validating epoch 9

Epoch 9 val loss: 1.6561
Saving new best model..
Training epoch 10...
Current LR: 0.002

Epoch 10 train loss: 1.9183
Validating epoch 10

Epoch 10 val loss: 1.6562

Run history:



Run summary:

lr	0.002
train_loss	1.91827
val_loss	1.65616

View run **chocolate-sunset-10** at: <https://wandb.ai/demonstem/precipitation-nowcasting/runs/4q5utifx> (<https://wandb.ai/demonstem/precipitation-nowcasting/runs/4q5utifx>)
Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

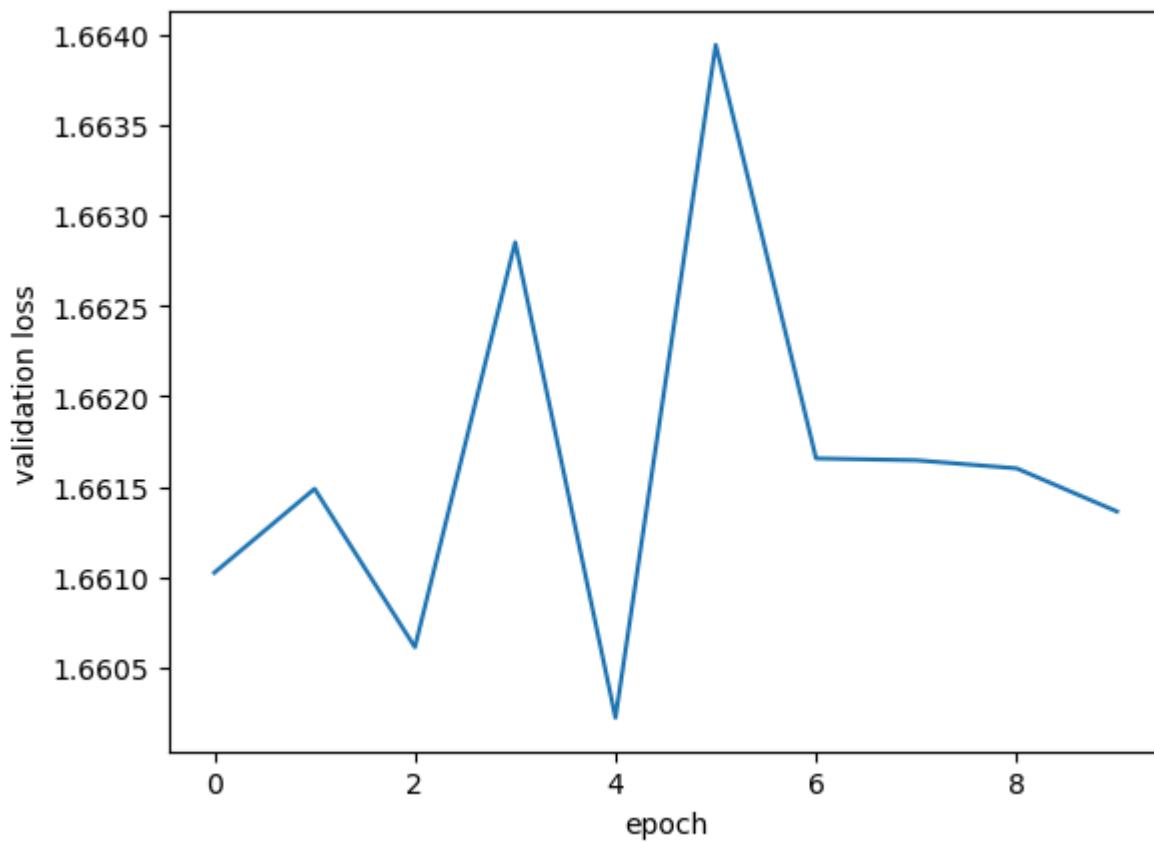
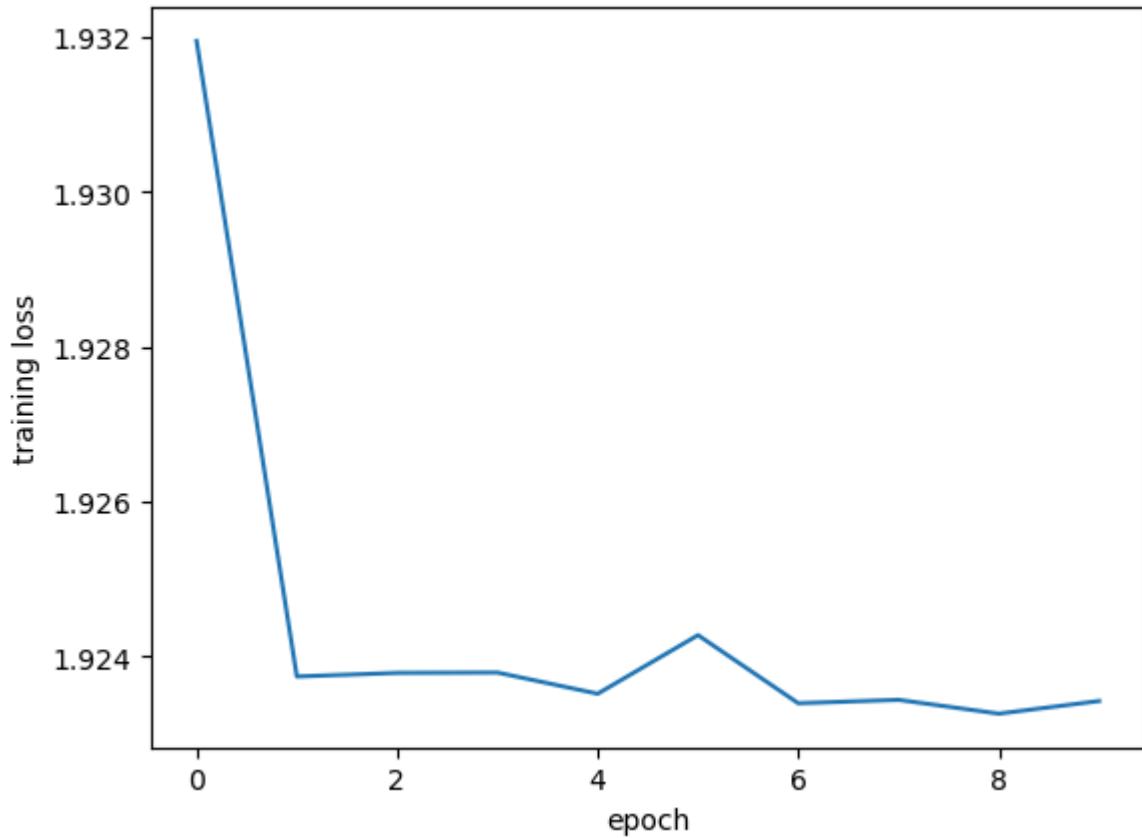
Find logs at: .\wandb\run-20240319_235818-4q5utifx\logs

Finished Training

TODO#5

Plot loss and val_loss as a function of epochs.

```
In [17]: plt.plot(train_losses)
plt.xlabel('epoch')
plt.ylabel('training loss')
plt.show()
plt.plot(val_losses)
plt.xlabel('epoch')
plt.ylabel('validation loss')
plt.show()
```



TODO#6

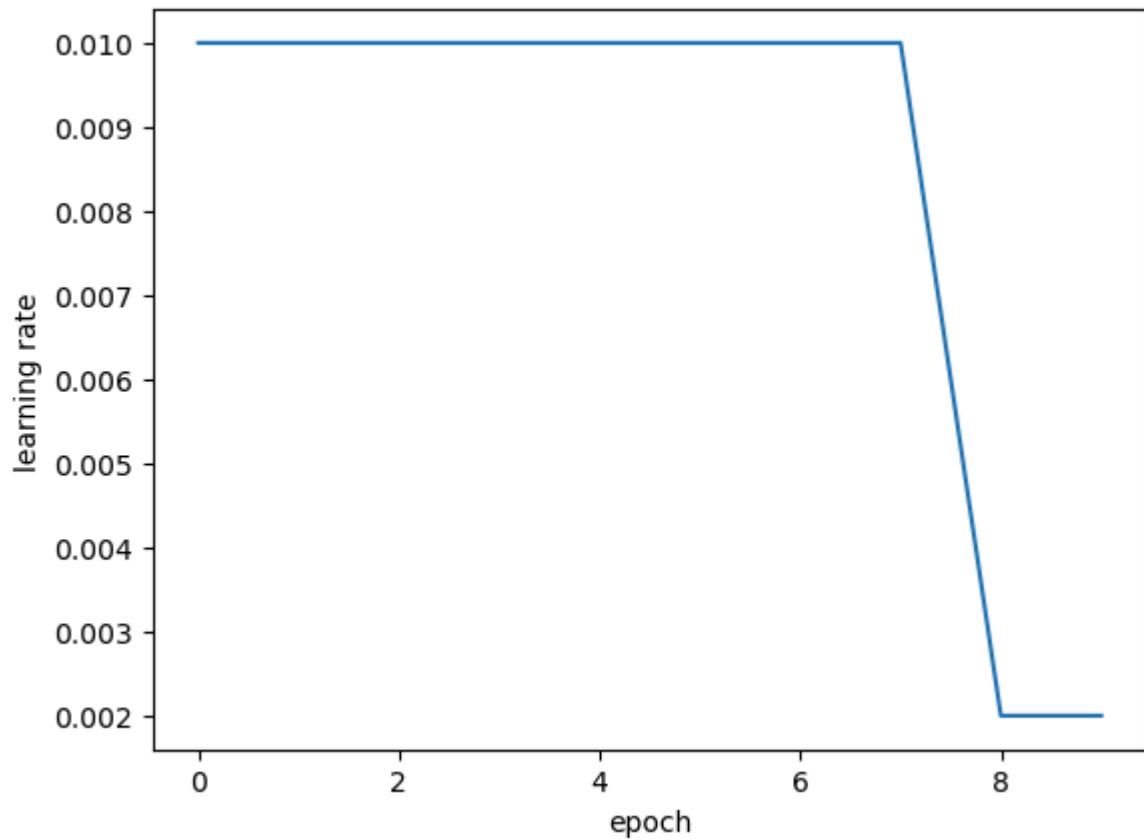
When does the model start to overfit?

Ans: at epoch 1 (observe from both train loss and val loss).

TODO#7

Plot the learning rate as a function of the epochs.

```
In [18]: plt.plot(learning_rates)
plt.xlabel('epoch')
plt.ylabel('learning rate')
plt.show()
```



TODO#8

What makes the learning rate change? (hint: try to understand the scheduler [ReduceLROnPlateau](https://pytorch.org/docs/stable/generated/torch.optim.lr_scheduler.ReduceLROnPlateau.html) (https://pytorch.org/docs/stable/generated/torch.optim.lr_scheduler.ReduceLROnPlateau.html))

Ans: learning rate is reduced when val_loss is not getting better.

Load Model

Use the code snippet below to load the model you just trained

```
In [15]: checkpoint = torch.load('model_ff.pth.tar')
loaded_model = FeedForwardNN(hidden_size=config['hidden_size']) # Create model
object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
```

Loaded epoch 4 model

```
In [16]: model_ff = loaded_model
```

A more complex scheduling

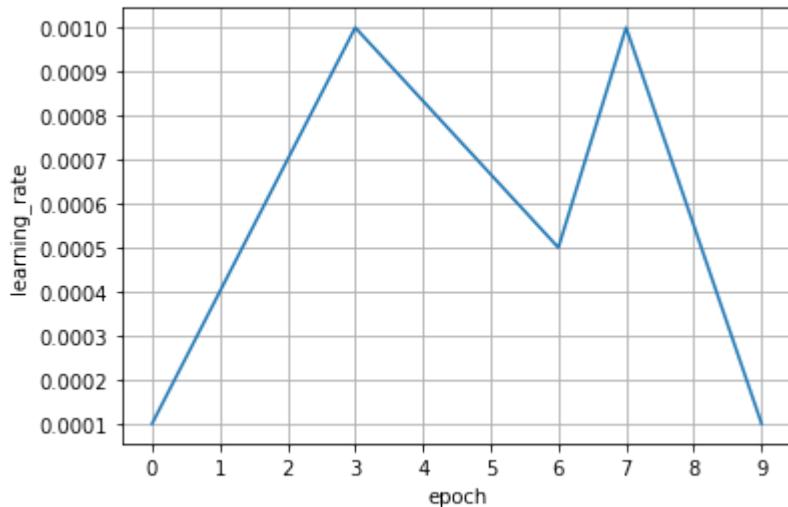
The scheduler can be very complicated and you can write your own heuristic for it.

TODO#9

Implement a custom learning rate scheduler that behaves like the following graph.

You might want to learn how to use [PyTorch's built-in learning rate schedulers](#) (<https://pytorch.org/docs/stable/optim.html#how-to-adjust-learning-rate>) in order to build your own.

Learning rate should be function of epoch.



In [17]: # Implement scheduler here

```
class MyScheduler():
    def __init__(self, optimizer: torch.optim.Optimizer):
        self.slope = [0.0003, 0.0003, 0.0003, -0.0005/3, -0.0005/3, -0.0005/3,
0.0005, -0.0009/2, -0.0009/2]
        self.lr = 0.0001
        self.optimizer = optimizer
        # for param_group in self.optimizer.param_groups:
        #     param_group['lr'] = self.lr
    def step(self, epoch):
        if epoch >= len(self.slope):
            print('!exceed 10 epoch!')
        else:
            self.lr += self.slope[epoch]
            for param_group in self.optimizer.param_groups:
                param_group['lr'] = self.lr
```

```
In [18]: # Now train with your scheduler
# my_scheduler = MyScheduler(...)

custom_scheduler_config = {
    'architecture': 'feedforward',
    'lr': 0.0001,
    'hidden_size': 200,
    'scheduler_factor': 0.2,
    'scheduler_patience': 2,
    'scheduler_min_lr': 1e-4,
    'epochs': 10
}

model_custom_scheduler = FeedForwardNN(hidden_size=custom_scheduler_config['hidden_size'])
model_custom_scheduler = model_custom_scheduler.to(device)

optimizer = torch.optim.Adam(model_custom_scheduler.parameters(), lr=custom_scheduler_config['lr'])
my_scheduler = MyScheduler(optimizer)

train_losses_cs = []
val_losses_cs = []
learning_rates_cs = []

# Start wandb run
wandb.init(
    project='precipitation-nowcasting-custom-scheduler',
    config=custom_scheduler_config,
)

# Log parameters and gradients
wandb.watch(model_custom_scheduler, log='all')

for epoch in range(custom_scheduler_config['epochs']): # Loop over the database multiple times

    # Training
    train_loss = []
    current_lr = optimizer.param_groups[0]['lr']
    learning_rates_cs.append(current_lr)

    # Flag model as training. Some layers behave differently in training and
    # inference modes, such as dropout, BN, etc.
    model_custom_scheduler.train()

    print(f"Training epoch {epoch+1}...")
    print(f"Current LR: {current_lr}")

    for i, (inputs, y_true) in enumerate(tqdm(train_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Reset the gradient
        optimizer.zero_grad()
```

```

# Predict
y_pred = model_custom_scheduler(inputs)

# Calculate loss
loss = loss_fn(y_pred, y_true)

# Compute gradient
loss.backward()

# Update parameters
optimizer.step()

# Log stuff
train_loss.append(loss)

avg_train_loss = torch.stack(train_loss).mean().item()
train_losses_cs.append(avg_train_loss)

print(f"Epoch {epoch+1} train loss: {avg_train_loss:.4f}")

# Validation
model_custom_scheduler.eval()
with torch.no_grad(): # No gradient is required during validation
    print(f"Validating epoch {epoch+1}")
    val_loss = []
    for i, (inputs, y_true) in enumerate(tqdm(val_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Predict
        y_pred = model_custom_scheduler(inputs)

        # Calculate loss
        loss = loss_fn(y_pred, y_true)

        # Log stuff
        val_loss.append(loss)

    avg_val_loss = torch.stack(val_loss).mean().item()
    val_losses_cs.append(avg_val_loss)
    print(f"Epoch {epoch+1} val loss: {avg_val_loss:.4f}")

# LR adjustment with scheduler
my_scheduler.step(epoch)

# Save checkpoint if val_loss is the best we got
best_val_loss = np.inf if epoch == 0 else min(val_losses_cs[:-1])
if avg_val_loss < best_val_loss:
    # Save whatever you want
    state = {
        'epoch': epoch,
        'model': model_custom_scheduler.state_dict(),
        'optimizer': optimizer.state_dict(),
        'scheduler': scheduler.state_dict(),
        'train_loss': avg_train_loss,

```

```
        'val_loss': avg_val_loss,
        'best_val_loss': best_val_loss,
    }

    print(f"Saving new best model..")
    torch.save(state, 'model_custom_scheduler.pth.tar')

wandb.log({
    'train_loss': avg_train_loss,
    'val_loss': avg_val_loss,
    'lr': current_lr,
})

wandb.finish()
print('Finished Training')
```

Failed to detect the name of this notebook, you can set it manually with the WANDB_NOTEBOOK_NAME environment variable to enable code saving.

wandb: Currently logged in as: `demonstem`. Use `wandb login --relogin` to force relogin

Tracking run with wandb version 0.16.4

Run data is saved locally in `c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240318_230820-0ly9qv4s`

Syncing run [lilac-monkey-6 \(https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler/runs/0ly9qv4s\)](https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler/runs/0ly9qv4s) to Weights & Biases (<https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler>) ([docs](https://wandb.me/run) (<https://wandb.me/run>))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler> (<https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler/runs/0ly9qv4s> (<https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler/runs/0ly9qv4s>)

Training epoch 1...
Current LR: 0.0001

Epoch 1 train loss: 1.9192
Validating epoch 1

Epoch 1 val loss: 1.6572
Saving new best model..
Training epoch 2...
Current LR: 0.00039999999999999996

Epoch 2 train loss: 1.9189
Validating epoch 2

Epoch 2 val loss: 1.6579
Training epoch 3...
Current LR: 0.0006999999999999999

Epoch 3 train loss: 1.9183
Validating epoch 3

Epoch 3 val loss: 1.6563
Saving new best model..
Training epoch 4...
Current LR: 0.0009999999999999998

Epoch 4 train loss: 1.9192
Validating epoch 4

Epoch 4 val loss: 1.6587
Training epoch 5...
Current LR: 0.00083333333333332

Epoch 5 train loss: 1.9184
Validating epoch 5

Epoch 5 val loss: 1.6565
Training epoch 6...
Current LR: 0.0006666666666666665

Epoch 6 train loss: 1.9179
Validating epoch 6

Epoch 6 val loss: 1.6560
Saving new best model..
Training epoch 7...
Current LR: 0.0004999999999999999

Epoch 7 train loss: 1.9176
Validating epoch 7

Epoch 7 val loss: 1.6563
Training epoch 8...
Current LR: 0.001

Epoch 8 train loss: 1.9183
Validating epoch 8

Epoch 8 val loss: 1.6569
Training epoch 9...
Current LR: 0.00055

Epoch 9 train loss: 1.9175
Validating epoch 9

Epoch 9 val loss: 1.6566
Training epoch 10...
Current LR: 0.0001000000000000005

Epoch 10 train loss: 1.9175
Validating epoch 10

Epoch 10 val loss: 1.6557
!exceed 10 epoch!
Saving new best model..

Run history:



Run summary:

lr	0.0001
train_loss	1.9175
val_loss	1.6557

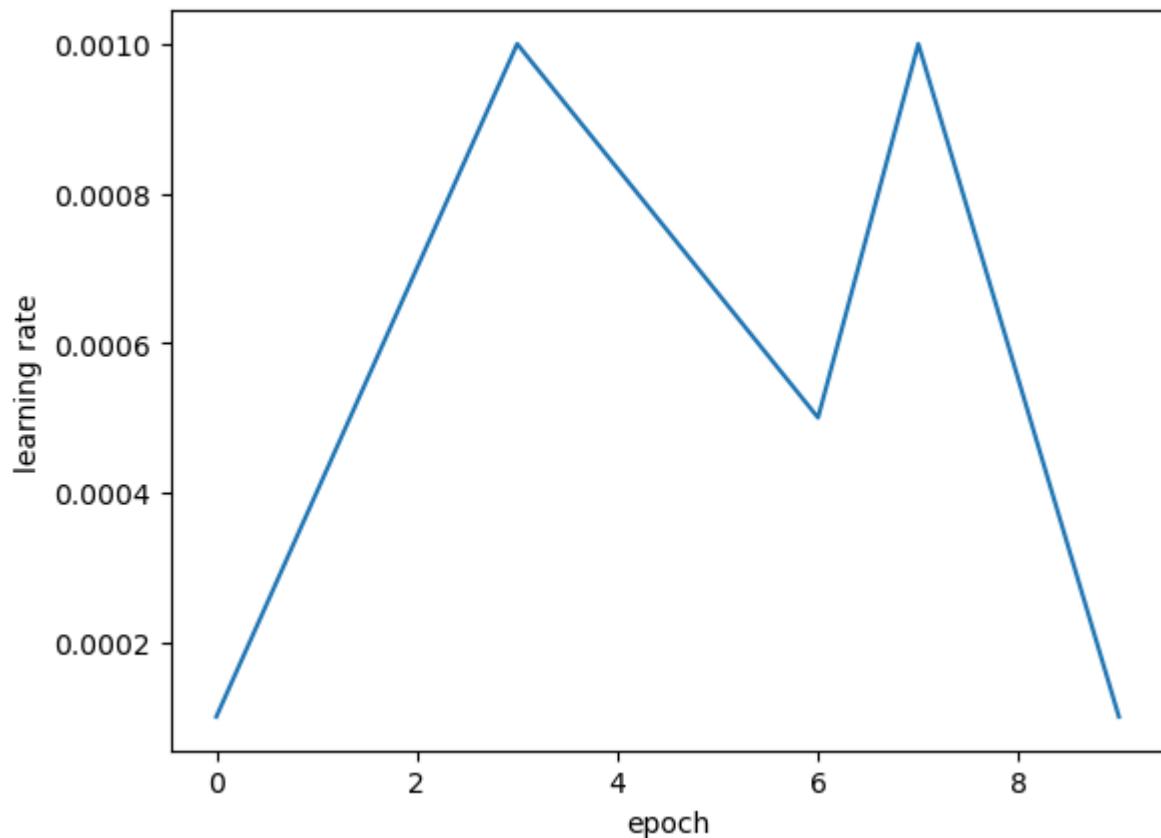
View run **lilac-monkey-6** at: <https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler/runs/0ly9qv4s> (<https://wandb.ai/demonstem/precipitation-nowcasting-custom-scheduler/runs/0ly9qv4s>)

Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240318_230820-0ly9qv4s\logs

Finished Training

```
In [19]: plt.plot(learning_rates_cs)
plt.xlabel('epoch')
plt.ylabel('learning rate')
plt.show()
```



[Optional] Wandb

You should now have a project in wandb with the name `precipitation-nowcasting`, which you should see the latest run you just finished inside the project. If you look into the run, you should be able to see plots of learning rate, train loss, val loss in the `Charts` section. Below it should be `Gradients` and `Parameters` section.

Wandb Observation

Optional TODO#1

Write your own interpretation of the logs from this example. A simple sentence or two for each section is sufficient.

Your answer: it overfitted at epoch 1 since the training loss did not increased but the validation loss is increased.

Evaluation

```
In [71]: #####
## 
# TODO#10:
#
# Write a function to evaluate your model. Your function must predicts
#
# using the input model and return mean square error of the model.
#
#
#
# Hint: Read how to use PyTorch's MSE Loss
#
# https://pytorch.org/docs/stable/generated/torch.nn.MSELoss.html
#
#####
##                                     WRITE YOUR CODE BELOW
#
#####
##
```

`#`

`def evaluate(data_loader, model):`

`"""`

`Evaluate model on validation data given by data_loader`

`"""`

`# write code here`

`losses = []`

`for inputs, y_true in tqdm(data_loader):`

`inputs, y_true = inputs.to(device), y_true.to(device)`

`y_pred = model(inputs)`

`losses.append(loss_fn(y_pred, y_true))`

`return sum(losses) / len(losses)`

```
In [21]: # We will use majority rule as a baseline.
def majority_baseline(label_set):
    unique, counts = np.unique(label_set, return_counts=True)
    majority = unique[np.argmax(counts)]
    baseline = 0
    label_set = label_set.reshape(-1,1)
    for r in label_set:
        baseline += (majority - r) ** 2 / len(label_set)
    return baseline
```

```
In [47]: print('baseline')
print('train', majority_baseline(y_train))
print('validate', majority_baseline(y_val))
```

```
baseline
train [1.94397725]
validate [1.6746546]
```

```
In [24]: print('FF-model')
model_ff = model_ff.to(device)
print('train', evaluate(train_loader, model_ff).item())
print('validate', evaluate(val_loader, model_ff).item())
```

FF-model

train 1.9241724014282227

validate 1.6602238416671753

Dropout

You might notice that the 3-layered feedforward does not use dropout at all. Now, try adding dropout (dropout rate of 20%) to the model, run, and report the result again.

To access PyTorch's dropout, use `nn.Dropout`. Read more about PyTorch's built-in Dropout layer [here](https://pytorch.org/docs/stable/generated/torch.nn.Dropout.html) (<https://pytorch.org/docs/stable/generated/torch.nn.Dropout.html>)

```
In [41]: #####
##
# TODO#11:
#
# Write a feedforward model with dropout
#
#####
##                                     WRITE YOUR CODE BELOW
#
#####
#
class DropOutNN(nn.Module):
    def __init__(self, hidden_size=200):
        super(DropOutNN, self).__init__()
        self.ff1 = nn.Linear(75, hidden_size)
        self.ff2 = nn.Linear(hidden_size, hidden_size)
        self.ff3 = nn.Linear(hidden_size, hidden_size)
        self.out = nn.Linear(hidden_size, 1)
        self.dropout = nn.Dropout(0.2)

    def forward(self, x):
        x = F.relu(self.ff1(x))
        x = self.dropout(x)
        x = F.relu(self.ff2(x))
        x = self.dropout(x)
        x = F.relu(self.ff3(x))
        x = self.dropout(x)
        x = self.out(x)
        return x.reshape(-1, 1)
```

```
In [42]: #####  
##  
# TODO#12:  
#  
# Complete the code to train your dropout model  
#  
#####  
##  
print('start training ff dropout')  
#####  
##  
#  
#           WRITE YOUR CODE BELOW  
#  
#####  
##  
# Hyperparameters and other configs  
dropout_config = {  
    'architecture': 'feedforward',  
    'lr': 0.01,  
    'hidden_size': 200,  
    'scheduler_factor': 0.2,  
    'scheduler_patience': 2,  
    'scheduler_min_lr': 1e-4,  
    'epochs': 10  
}  
  
# Model  
model_dropout = DropOutNN(hidden_size=dropout_config['hidden_size'])  
model_dropout = model_dropout.to(device)  
optimizer_dropout = torch.optim.Adam(model_dropout.parameters(), lr=dropout_co  
nfig['lr'])  
scheduler_dropout = torch.optim.lr_scheduler.ReduceLROnPlateau(  
    optimizer_dropout,  
    'min',  
    factor=dropout_config['scheduler_factor'],  
    patience=dropout_config['scheduler_patience'],  
    min_lr=dropout_config['scheduler_min_lr'])  
)  
  
# Training  
train_losses_dropout = []  
val_losses_dropout = []  
learning_rates_dropout = []  
  
# Start wandb run  
wandb.init(  
    project='precipitation-nowcasting-dropout',  
    config=dropout_config,  
)  
  
# Log parameters and gradients  
wandb.watch(model_dropout, log='all')  
  
for epoch in range(dropout_config['epochs']): # Loop over the dataset multipl  
e times
```

```

# Training
train_loss = []
current_lr = optimizer_dropout.param_groups[0]['lr']
learning_rates_dropout.append(current_lr)

# Flag model as training. Some layers behave differently in training and
# inference modes, such as dropout, BN, etc.
model_dropout.train()

print(f"Training epoch {epoch+1}...")
print(f"Current LR: {current_lr}")

for i, (inputs, y_true) in enumerate(tqdm(train_loader)):
    # Transfer data from cpu to gpu
    inputs = inputs.to(device)
    y_true = y_true.to(device)

    # Reset the gradient
    optimizer_dropout.zero_grad()

    # Predict
    y_pred = model_dropout(inputs)

    # Calculate loss
    loss = loss_fn(y_pred, y_true)

    # Compute gradient
    loss.backward()

    # Update parameters
    optimizer_dropout.step()

    # Log stuff
    train_loss.append(loss)

avg_train_loss = torch.stack(train_loss).mean().item()
train_losses_dropout.append(avg_train_loss)

print(f"Epoch {epoch+1} train loss: {avg_train_loss:.4f}")

# Validation
model_dropout.eval()
with torch.no_grad(): # No gradient is required during validation
    print(f"Validating epoch {epoch+1}")
    val_loss = []
    for i, (inputs, y_true) in enumerate(tqdm(val_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Predict
        y_pred = model_dropout(inputs)

        # Calculate loss
        loss = loss_fn(y_pred, y_true)

        # Log stuff

```

```
    val_loss.append(loss)

    avg_val_loss = torch.stack(val_loss).mean().item()
    val_losses_dropout.append(avg_val_loss)
    print(f"Epoch {epoch+1} val loss: {avg_val_loss:.4f}")

    # LR adjustment with scheduler
    scheduler_dropout.step(epoch)

    # Save checkpoint if val_loss is the best we got
    best_val_loss = np.inf if epoch == 0 else min(val_losses_dropout[:-1])
    if avg_val_loss < best_val_loss:
        # Save whatever you want
        state = {
            'epoch': epoch,
            'model': model_dropout.state_dict(),
            'optimizer': optimizer_dropout.state_dict(),
            'scheduler': scheduler.state_dict(),
            'train_loss': avg_train_loss,
            'val_loss': avg_val_loss,
            'best_val_loss': best_val_loss,
        }

        print(f"Saving new best model..")
        torch.save(state, 'model_dropout.pth.tar')

    wandb.log({
        'train_loss': avg_train_loss,
        'val_loss': avg_val_loss,
        'lr': current_lr,
    })

wandb.finish()
print('Finished Training')
```

start training ff dropout

Tracking run with wandb version 0.16.4

Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240319_015528-cfyc2qhk

Syncing run [wandering-sound-1 \(https://wandb.ai/demonstem/precipitation-nowcasting-dropout/runs/cfyc2qhk\)](https://wandb.ai/demonstem/precipitation-nowcasting-dropout/runs/cfyc2qhk) to [Weights & Biases \(https://wandb.ai/demonstem/precipitation-nowcasting-dropout\)](https://wandb.ai/demonstem/precipitation-nowcasting-dropout) ([docs \(https://wandb.me/run\)](https://wandb.me/run))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-dropout> (<https://wandb.ai/demonstem/precipitation-nowcasting-dropout>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-dropout/runs/cfyc2qhk> (<https://wandb.ai/demonstem/precipitation-nowcasting-dropout/runs/cfyc2qhk>)

Training epoch 1...
Current LR: 0.01

Epoch 1 train loss: 1.9359
Validating epoch 1

Epoch 1 val loss: 1.6621
Saving new best model..
Training epoch 2...
Current LR: 0.01

Epoch 2 train loss: 1.9233
Validating epoch 2

Epoch 2 val loss: 1.6605
Saving new best model..
Training epoch 3...
Current LR: 0.01

Epoch 3 train loss: 1.9237
Validating epoch 3

Epoch 3 val loss: 1.6610
Training epoch 4...
Current LR: 0.01

Epoch 4 train loss: 1.9235
Validating epoch 4

Epoch 4 val loss: 1.6606
Training epoch 5...
Current LR: 0.002

Epoch 5 train loss: 1.9235
Validating epoch 5

Epoch 5 val loss: 1.6627

Training epoch 6...

Current LR: 0.002

Epoch 6 train loss: 1.9240

Validating epoch 6

Epoch 6 val loss: 1.6609

Training epoch 7...

Current LR: 0.002

Epoch 7 train loss: 1.9233

Validating epoch 7

Epoch 7 val loss: 1.6610

Training epoch 8...

Current LR: 0.0004

Epoch 8 train loss: 1.9236

Validating epoch 8

Epoch 8 val loss: 1.6616

Training epoch 9...

Current LR: 0.0004

Epoch 9 train loss: 1.9233

Validating epoch 9

Epoch 9 val loss: 1.6610

Training epoch 10...

Current LR: 0.0004

Epoch 10 train loss: 1.9233

Validating epoch 10

Epoch 10 val loss: 1.6613

Run history:



Run summary:

lr	0.0004
train_loss	1.92333
val_loss	1.66125

View run **wandering-sound-1** at: <https://wandb.ai/demonstem/precipitation-nowcasting-dropout/runs/cfyc2qhk> (<https://wandb.ai/demonstem/precipitation-nowcasting-dropout/runs/cfyc2qhk>)

Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

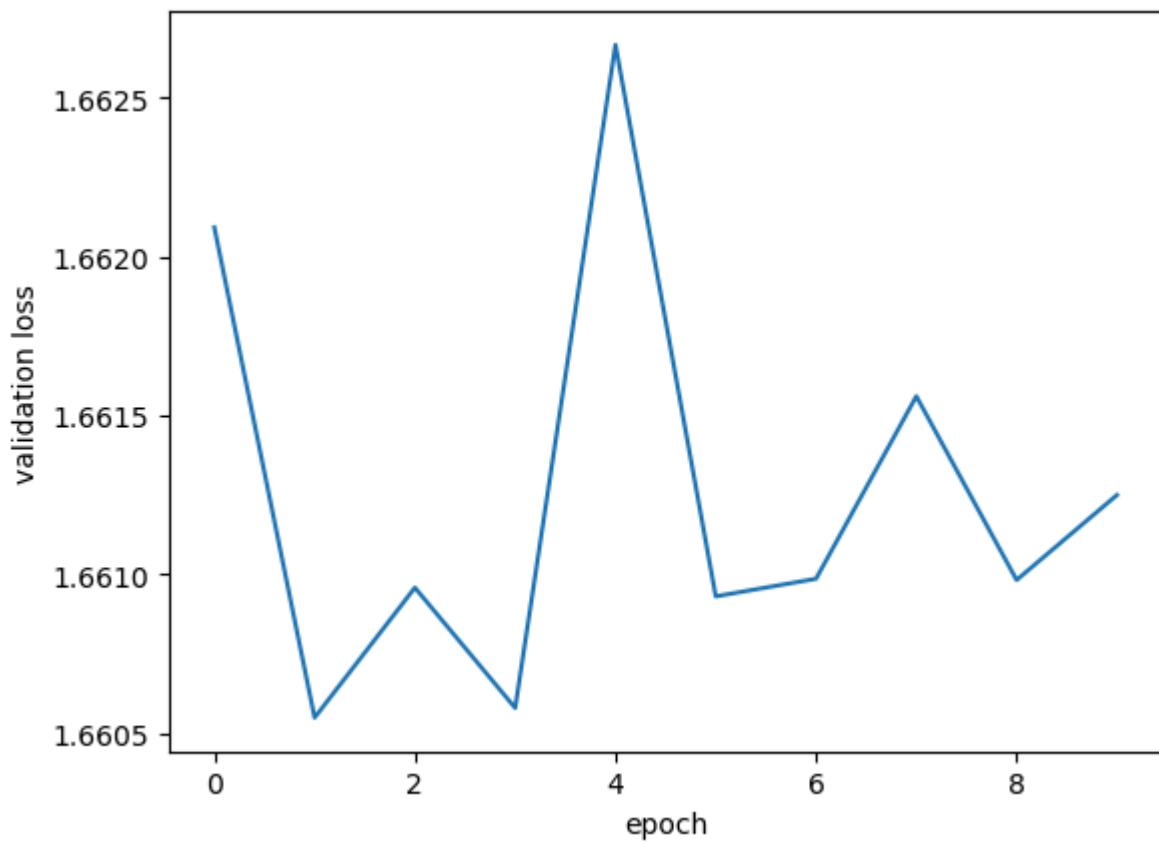
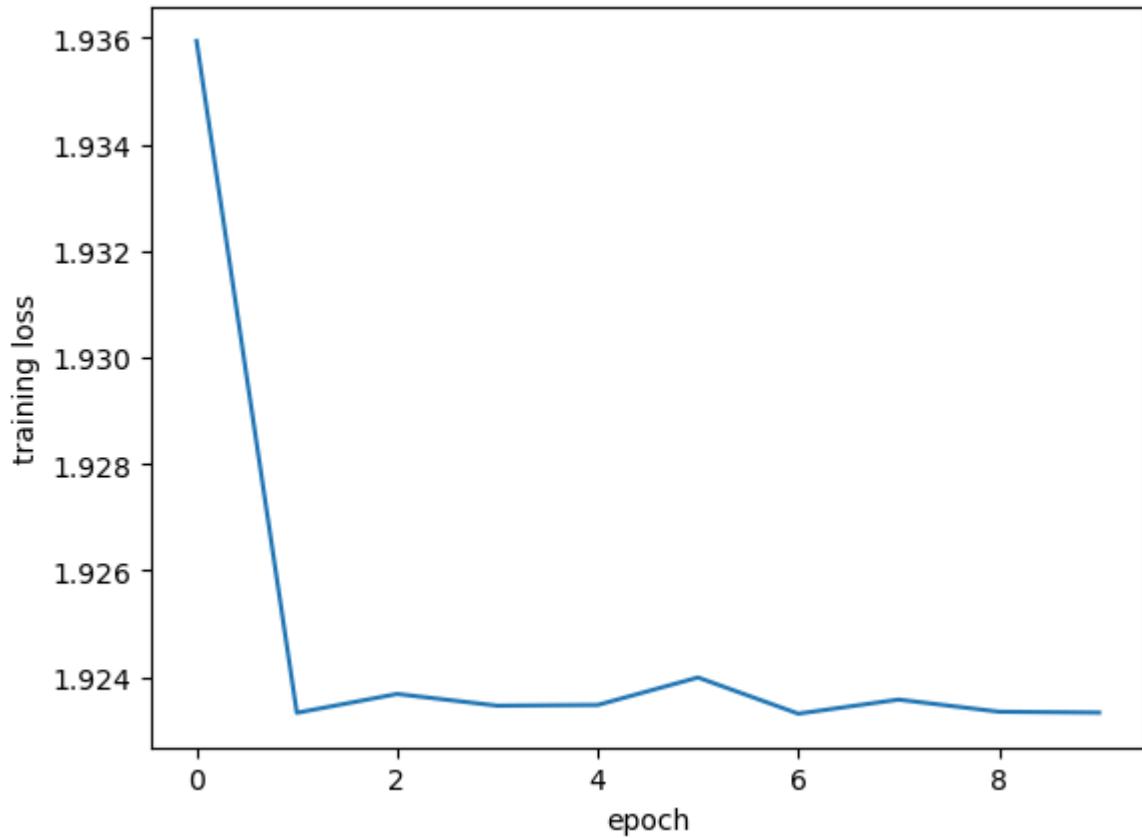
Find logs at: .\wandb\run-20240319_015528-cfyc2qhk\logs

Finished Training

TODO#13

Plot the losses and MSE of the training and validation as before. Evaluate the dropout model's performance

```
In [64]: # Plot here
plt.plot(train_losses_dropout)
plt.xlabel('epoch')
plt.ylabel('training loss')
plt.show()
plt.plot(val_losses_dropout)
plt.xlabel('epoch')
plt.ylabel('validation loss')
plt.show()
```



```
In [63]: # Evaluate
print('Dropout-model')
print('train', evaluate(train_loader, model_dropout).item())
print('validate', evaluate(val_loader, model_dropout).item())
```

```
Dropout-model
```

```
train 1.92324960231781
```

```
validate 1.6612495183944702
```

Convolution Neural Networks

Now let's try to incorporate the grid structure to your model. Instead of passing in vectors, we are going to pass in the 5x5 grid into the model (5lat x 5long x 3channel). You are going to implement your own 2d-convolution neural networks with the following structure.

```
=====
=====
Layer (type:depth-idx)          Output Shape      Param #
=====
=====
Conv2DNN                         --                  --
|Conv2d: 1-1                      [1024, 200, 3, 3]    5,600
|Linear: 1-2                       [1024, 200]        360,200
|Linear: 1-3                       [1024, 200]        40,200
|Linear: 1-4                       [1024, 1]          201
=====
=====
Total params: 406,201
Trainable params: 406,201
Non-trainable params: 0
```

These parameters are simple guidelines to save your time.

You can play with them in the final section which you can choose any normalization methods, activation function, as well as any hyperparameter the way you want.

Hint: You should read PyTorch documentation to see the list of available layers and options you can use.

In [63]:

```
#####
##  
# TODO#14:  
#  
# Complete the code for preparing data for training CNN  
#  
# Input for CNN should not have time step.  
#  
#####  
##  
#                      WRITE YOUR CODE BELOW  
#  
#####  
##  
def to_cnn_inputs(x):  
    return torch.reshape(x, (-1, 3, 5, 5))
```

```
In [64]: #####  
##  
# TODO#15:  
#  
# Write a PyTorch convolutional neural network model.  
#  
# You might want to use the layer torch.flatten somewhere  
#  
#####  
##  
#          WRITE YOUR CODE BELOW  
#  
#####  
##  
class CNN(nn.Module):  
    def __init__(self, hidden_size=200):  
        super(CNN, self).__init__()  
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=hidden_size, kernel_size=3)  
        self.ff1 = nn.Linear(hidden_size * 3 * 3, hidden_size)  
        self.ff2 = nn.Linear(hidden_size, hidden_size)  
        self.ff3 = nn.Linear(hidden_size, 1)  
  
    def forward(self, x):  
        x = torch.relu(self.conv1(x))  
        x = torch.flatten(x, 1)  
        x = torch.relu(self.ff1(x))  
        x = torch.relu(self.ff2(x))  
        x = self.ff3(x)  
        return x  
  
# Hyperparameters and configs  
config_cnn = {  
    'architecture': 'cnn',  
    'lr': 0.01,  
    'hidden_size': 200,  
    'scheduler_factor': 0.2,  
    'scheduler_patience': 2,  
    'scheduler_min_lr': 1e-4,  
    'epochs': 10,  
}  
# Model  
model_cnn = CNN(hidden_size=config_cnn['hidden_size'])  
model_cnn = model_cnn.to(device)  
summary(model_cnn, input_size=(1024, 3, 5, 5))
```

```
Out[64]: =====
=====
Layer (type:depth-idx)          Output Shape      Param #
=====
=====
CNN                               [1024, 1]           --
|Conv2d: 1-1                     [1024, 200, 3, 3]    5,600
|Linear: 1-2                      [1024, 200]         360,200
|Linear: 1-3                      [1024, 200]         40,200
|Linear: 1-4                      [1024, 1]           201
=====
=====
Total params: 406,201
Trainable params: 406,201
Non-trainable params: 0
Total mult-adds (Units.MEGABYTES): 461.83
=====
=====
Input size (MB): 0.31
Forward/backward pass size (MB): 18.03
Params size (MB): 1.62
Estimated Total Size (MB): 19.96
=====
```

In [21]:

```
#####
## 
# TODO#16:
#
# Complete the code to train your cnn model
#
#####
## 
print('start training conv2d')
#####
## 
##           WRITE YOUR CODE BELOW
#
#####
## 
optimizer_cnn = torch.optim.Adam(model_cnn.parameters(), lr=config_cnn['lr'])
scheduler_cnn = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer_cnn,
    'min',
    factor=config_cnn['scheduler_factor'],
    patience=config_cnn['scheduler_patience'],
    min_lr=config_cnn['scheduler_min_lr']
)

train_losses_cnn = []
val_losses_cnn = []
learning_rates_cnn = []

# Start wandb run
wandb.init(
    project='precipitation-nowcasting-cnn',
    config=config_cnn,
)
# Log parameters and gradients
wandb.watch(model_cnn, log='all')

for epoch in range(config_cnn['epochs']): # Loop over the dataset multiple times

    # Training
    train_loss = []
    current_lr = optimizer_cnn.param_groups[0]['lr']
    learning_rates_cnn.append(current_lr)

    # Flag model as training. Some layers behave differently in training and
    # inference modes, such as dropout, BN, etc.
    model_cnn.train()

    print(f"Training epoch {epoch+1}...")
    print(f"Current LR: {current_lr}")

    for i, (inputs, y_true) in enumerate(tqdm(train_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)
```

```

# Reset the gradient
optimizer_cnn.zero_grad()

# Predict
y_pred = model_cnn(to_cnn_inputs(inputs))

# Calculate loss
loss = loss_fn(y_pred, y_true)

# Compute gradient
loss.backward()

# Update parameters
optimizer_cnn.step()

# Log stuff
train_loss.append(loss)

avg_train_loss = torch.stack(train_loss).mean().item()
train_losses_cnn.append(avg_train_loss)

print(f"Epoch {epoch+1} train loss: {avg_train_loss:.4f}")

# Validation
model_cnn.eval()
with torch.no_grad(): # No gradient is required during validation
    print(f"Validating epoch {epoch+1}")
    val_loss = []
    for i, (inputs, y_true) in enumerate(tqdm(val_loader)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Predict
        y_pred = model_cnn(to_cnn_inputs(inputs))

        # Calculate loss
        loss = loss_fn(y_pred, y_true)

        # Log stuff
        val_loss.append(loss)

    avg_val_loss = torch.stack(val_loss).mean().item()
    val_losses_cnn.append(avg_val_loss)
    print(f"Epoch {epoch+1} val loss: {avg_val_loss:.4f}")

# LR adjustment with my_scheduler
scheduler_cnn.step(avg_val_loss)

# Save checkpoint if val_loss is the best we got
best_val_loss = np.inf if epoch == 0 else min(val_losses_cnn[:-1])
if avg_val_loss < best_val_loss:
    # Save whatever you want
    state = {
        'epoch': epoch,
        'model': model_cnn.state_dict(),

```

```
'optimizer': optimizer_cnn.state_dict(),
'scheduler': scheduler_cnn.state_dict(),
'train_loss': avg_train_loss,
'val_loss': avg_val_loss,
'best_val_loss': best_val_loss,
}

print(f"Saving new best model..")
torch.save(state, 'model_cnn.pth.tar')

wandb.log({
    'train_loss': avg_train_loss,
    'val_loss': avg_val_loss,
    'lr': current_lr,
})

wandb.finish()
print('Finished Training')
```

```
start training conv2d
```

Failed to detect the name of this notebook, you can set it manually with the WANDB_NOTEBOOK_NAME environment variable to enable code saving.
wandb: Currently logged in as: `demonstem`. Use `wandb login --relogin` to force relogin

Tracking run with wandb version 0.16.4

Run data is saved locally in `c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_002541-9uheewvx`

Syncing run [polished-puddle-2 \(https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/9uheewvx\)](https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/9uheewvx) to [Weights & Biases \(https://wandb.ai/demonstem/precipitation-nowcasting-cnn\)](https://wandb.ai/demonstem/precipitation-nowcasting-cnn) ([docs \(https://wandb.me/run\)](https://wandb.me/run))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/9uheewvx> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/9uheewvx>)

Training epoch 1...

Current LR: 0.01

Epoch 1 train loss: 2.1246

Validating epoch 1

Epoch 1 val loss: 1.6603

Saving new best model..

Training epoch 2...

Current LR: 0.01

Epoch 2 train loss: 1.9218

Validating epoch 2

Epoch 2 val loss: 1.6595

Saving new best model..

Training epoch 3...

Current LR: 0.01

Epoch 3 train loss: 1.9210

Validating epoch 3

Epoch 3 val loss: 1.6581

Saving new best model..

Training epoch 4...

Current LR: 0.01

Epoch 4 train loss: 1.9201

Validating epoch 4

Epoch 4 val loss: 1.6604

Training epoch 5...

Current LR: 0.01

Epoch 5 train loss: 1.9207

Validating epoch 5

Epoch 5 val loss: 1.6582

Training epoch 6...

Current LR: 0.01

Epoch 6 train loss: 1.9203

Validating epoch 6

Epoch 6 val loss: 1.6605

Training epoch 7...

Current LR: 0.002

Epoch 7 train loss: 1.9200

Validating epoch 7

Epoch 7 val loss: 1.6581

Training epoch 8...

Current LR: 0.002

Epoch 8 train loss: 1.9199

Validating epoch 8

Epoch 8 val loss: 1.6590

Training epoch 9...

Current LR: 0.002

Epoch 9 train loss: 1.9198

Validating epoch 9

Epoch 9 val loss: 1.6587

Training epoch 10...

Current LR: 0.0004

Epoch 10 train loss: 1.9196

Validating epoch 10

Epoch 10 val loss: 1.6583

Run history:



Run summary:

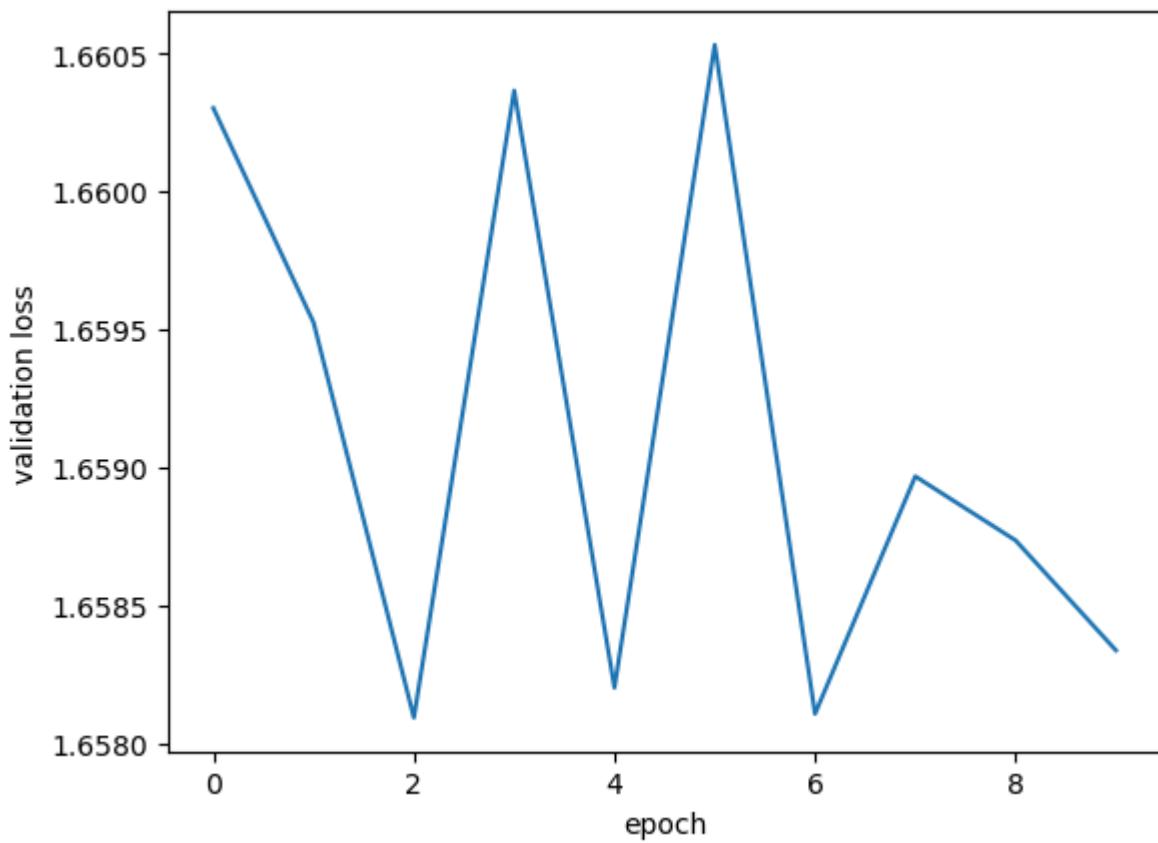
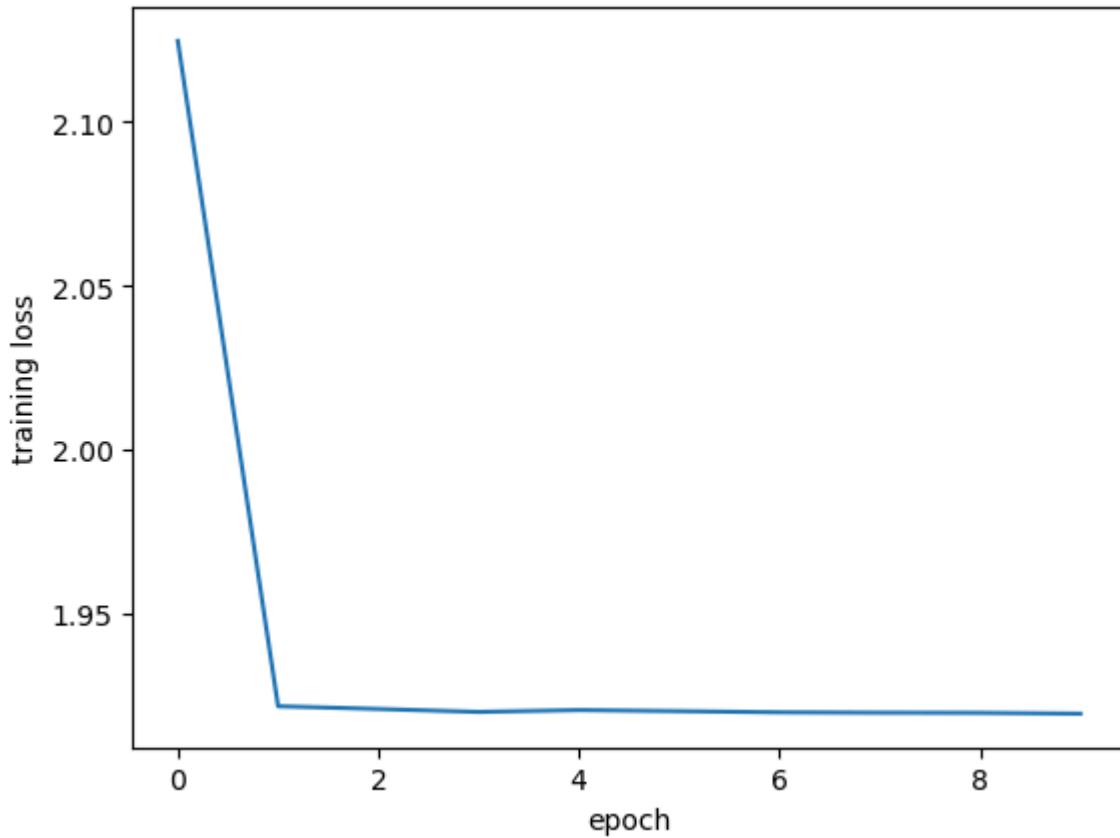
lr	0.0004
train_loss	1.91957
val_loss	1.65834

View run **polished-puddle-2** at: <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/9uheewvx> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/9uheewvx>)
Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_002541-9uheewvx\logs

Finished Training

```
In [22]: # Plot losses
plt.plot(train_losses_cnn)
plt.xlabel('epoch')
plt.ylabel('training loss')
plt.show()
plt.plot(val_losses_cnn)
plt.xlabel('epoch')
plt.ylabel('validation loss')
plt.show()
```



```
In [15]: checkpoint = torch.load('model_cnn.pth.tar')
loaded_model = CNN(hidden_size=config_cnn['hidden_size']) # Create model object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
model_cnn = loaded_model
model_cnn.to(device)
```

Loaded epoch 2 model

```
Out[15]: CNN(
    (conv1): Conv2d(3, 200, kernel_size=(3, 3), stride=(1, 1))
    (ff1): Linear(in_features=1800, out_features=200, bias=True)
    (ff2): Linear(in_features=200, out_features=200, bias=True)
    (ff3): Linear(in_features=200, out_features=1, bias=True)
)
```

```
In [65]: def evaluate_cnn(data_loader, model):
    wandb.init(
        project='precipitation-nowcasting-cnn',
    )

    # Log parameters and gradients
    wandb.watch(model, log='all')
    with torch.no_grad():
        train_losses = []
        for i, (inputs, y_true) in enumerate(tqdm(data_loader)):
            # Transfer data from cpu to gpu
            inputs = inputs.to(device)
            y_true = y_true.to(device)

            # Predict
            y_pred = model(to_cnn_inputs(inputs))

            # Calculate loss
            loss = loss_fn(y_pred, y_true)

            # Log stuff
            train_losses.append(loss)

        mse = torch.stack(train_losses).mean()
        wandb.finish()
    return mse
```

```
In [66]: # Evaluate
evaluated_train_cnn = evaluate_cnn(train_loader, model_cnn).item()
evaluated_val_cnn = evaluate_cnn(val_loader, model_cnn).item()

print('cnn-model')
print('train', evaluated_train_cnn)
print('validate', evaluated_val_cnn)
```

Finishing last run (ID:72op9xla) before initializing another...

View run **light-glitter-6** at: <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/72op9xla> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/72op9xla>)
Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_212520-72op9xla\logs

Successfully finished last run (ID:72op9xla). Initializing new run:

Tracking run with wandb version 0.16.4

Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_212641-uliy2zl5

Syncing run **hearty-firefly-7** (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/uliy2zl5>) to Weights & Biases (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>) ([docs \(https://wandb.me/run\)](https://wandb.me/run))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/uliy2zl5> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/uliy2zl5>)

View run **hearty-firefly-7** at: <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/uliy2zl5> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/uliy2zl5>)
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_212641-uliy2zl5\logs

Tracking run with wandb version 0.16.4

Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_212852-2r3l4brd

Syncing run **sparkling-vortex-8** (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/2r3l4brd>) to Weights & Biases (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>) ([docs \(https://wandb.me/run\)](https://wandb.me/run))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/2r3l4brd> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/2r3l4brd>)

View run **sparkling-vortex-8** at: <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/2r3l4brd> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/2r3l4brd>)
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_212852-2r3l4brd\logs

```
cnn-model
train 1.9381484985351562
validate 1.6678866147994995
```

Gated Recurrent Units

Now, you want to add time steps into your model. Recall the original data has 5 time steps per item. You are going to pass in a data of the form 5 timesteps x 75data. This can be done using a GRU layer. Implement your own GRU network with the following structure.

```
=====
=====
Layer (type:depth-idx)          Output Shape      Param #
=====
=====
GRUModel
├─GRU: 1-1                      [1024, 5, 200]    166,200
├─Linear: 1-2                    [1024, 5, 200]    40,200
└─Linear: 1-3                    [1024, 5, 1]     201
=====
=====
Total params: 206,601
Trainable params: 206,601
Non-trainable params: 0
```

These parameters are simple guidelines to save your time.

You can play with them in the final section which you can choose any normalization methods, activation function, as well as any hyperparameter the way you want.

The result should be better than the feedforward model and at least on par with your CNN model.

Do consult PyTorch documentation on how to use [GRUs](#)
(<https://pytorch.org/docs/stable/generated/torch.nn.GRU.html>).

In [24]:

```
#####
##  
# TODO#17:  
#  
# Complete the code for preparing data for training GRU  
#  
# GRU's input should has 3 dimensions.  
#  
# The dimensions should compose of entries, time-step, and features.  
#  
#####
##  
#  
#           WRITE YOUR CODE BELOW  
#  
#####
##  
def preprocess_for_gru(x_train, y_train, x_val, y_val):  
    x_train_gru = x_train.reshape((-1, 5, 5*5*3))  
    y_train_gru = y_train.reshape((-1, 5))  
    x_val_gru = x_val.reshape((-1, 5, 5*5*3))  
    y_val_gru = y_val.reshape((-1, 5))  
    x_test_gru = x_test.reshape((-1, 5, 5*5*3))  
    y_test_gru = y_test.reshape((-1, 5))  
  
    return x_train_gru, y_train_gru, x_val_gru, y_val_gru, x_test_gru, y_test_gru  
  
x_train_gru, y_train_gru, x_val_gru, y_val_gru, x_test_gru, y_test_gru = preprocess_for_gru(x_train, y_train, x_val, y_val)  
  
class RainfallDatasetGRU(Dataset):  
    def __init__(self, x, y, normalizer):  
        self.x = x.astype(np.float32)  
        self.y = y.astype(np.float32)  
        self.normalizer = normalizer  
        print(self.x.shape)  
        print(self.y.shape)  
  
    def __getitem__(self, index):  
        x = self.x[index] # Retrieve data  
        x = self.normalizer.transform(x.reshape(1, -1)) # Normalize  
        x = x.reshape(5, 75)  
        y = self.y[index]  
        return x, y  
  
    def __len__(self):  
        return self.x.shape[0]  
  
normalizer_gru = normalizer_std(x_train_gru.reshape(-1, 5*5*5*3)) # We will not  
# normalize everything based on x_train  
train_dataset_gru = RainfallDatasetGRU(x_train_gru, y_train_gru, normalizer_gru)  
val_dataset_gru = RainfallDatasetGRU(x_val_gru, y_val_gru, normalizer_gru)  
test_dataset_gru = RainfallDatasetGRU(x_test_gru, y_test_gru, normalizer_gru)
```

```

train_loader_gru = DataLoader(train_dataset_gru, batch_size=1024, shuffle=True,
                           pin_memory=True)
val_loader_gru = DataLoader(val_dataset_gru, batch_size=1024, shuffle=False,
                           pin_memory=True)
test_loader_gru = DataLoader(test_dataset_gru, batch_size=1024, shuffle=False,
                           pin_memory=True)

(229548, 5, 75)
(229548, 5)
(92839, 5, 75)
(92839, 5)
(111715, 5, 75)
(111715, 5)

```

In [56]:

```

#####
## 
# TODO#18
#
# Write a PyTorch GRU model.
#
# Your goal is to predict a precipitation of every time step.
#
#
# Hint: You should read PyTorch documentation to see the list of available
#
# Layers and options you can use.
#
#####
## 
#                               WRITE YOUR CODE BELOW
#
#####
## 
class GRU(nn.Module):
    def __init__(self, hidden_size=200):
        super(GRU, self).__init__()
        self.gru = nn.GRU(3 * 5 * 5, hidden_size, 5)
        self.ff1 = nn.Linear(hidden_size, hidden_size)
        self.out = nn.Linear(hidden_size, 1)

    def forward(self, x):
        gru = self.gru(x)
        hd1 = F.relu(self.ff1(gru[0]))
        y = self.out(hd1)
        return y

```

```
In [57]: config_gru = {
    'architecture': 'gru',
    'lr': 0.01,
    'hidden_size': 200,
    'scheduler_factor': 0.2,
    'scheduler_patience': 2,
    'scheduler_min_lr': 1e-4,
    'epochs': 10,
}

model_gru = GRU(hidden_size=config_gru['hidden_size'])
model_gru = model_gru.to(device)
optimizer_gru = torch.optim.Adam(model_gru.parameters(), lr=config_gru['lr'])
scheduler_gru = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer_gru,
    'min',
    factor=config_gru['scheduler_factor'],
    patience=config_gru['scheduler_patience'],
    min_lr=config_gru['scheduler_min_lr']
)
```

In [20]:

```
#####
## 
# TODO#19
#
# Complete the code to train your gru model
#
#####
## 
print('start training gru')
#####
## 
##           WRITE YOUR CODE BELOW
#
#####
## 
## 
train_losses_gru = []
val_losses_gru = []
learning_rates_gru = []

# Start wandb run
wandb.init(
    project='precipitation-nowcasting-gru',
    config=config_gru,
)

# Log parameters and gradients
wandb.watch(model_gru, log='all')

for epoch in range(config_gru['epochs']): # Loop over the dataset multiple times

    # Training
    train_loss = []
    current_lr = optimizer_gru.param_groups[0]['lr']
    learning_rates_gru.append(current_lr)

    # Flag model as training. Some Layers behave differently in training and
    # inference modes, such as dropout, BN, etc.
    model_gru.train()

    print(f"Training epoch {epoch+1}...")
    print(f"Current LR: {current_lr}")

    for i, (inputs, y_true) in enumerate(tqdm(train_loader_gru)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Reset the gradient
        optimizer_gru.zero_grad()

        # Predict
        y_pred = model_gru(inputs)

        # Calculate loss
        y_pred = torch.reshape(y_pred, (-1, 5))
```

```

        loss = loss_fn(y_pred, y_true)

        # Compute gradient
        loss.backward()

        # Update parameters
        optimizer_gru.step()

        # Log stuff
        train_loss.append(loss)

    avg_train_loss = torch.stack(train_loss).mean().item()
    train_losses_gru.append(avg_train_loss)

    print(f"Epoch {epoch+1} train loss: {avg_train_loss:.4f}")

    # Validation
    model_gru.eval()
    with torch.no_grad(): # No gradient is required during validation
        print(f"Validating epoch {epoch+1}")
        val_loss = []
        for i, (inputs, y_true) in enumerate(tqdm(val_loader_gru)):
            # Transfer data from cpu to gpu
            inputs = inputs.to(device)
            y_true = y_true.to(device)

            # Predict
            y_pred = model_gru(inputs)

            # Calculate loss
            y_pred = torch.reshape(y_pred, (-1, 5))
            loss = loss_fn(y_pred, y_true)

            # Log stuff
            val_loss.append(loss)

    avg_val_loss = torch.stack(val_loss).mean().item()
    val_losses_gru.append(avg_val_loss)
    print(f"Epoch {epoch+1} val loss: {avg_val_loss:.4f}")

    # LR adjustment with my_scheduler
    scheduler_gru.step(avg_val_loss)

    # Save checkpoint if val_loss is the best we got
    best_val_loss = np.inf if epoch == 0 else min(val_losses_gru[:-1])
    if avg_val_loss < best_val_loss:
        # Save whatever you want
        state = {
            'epoch': epoch,
            'model': model_gru.state_dict(),
            'optimizer': optimizer_gru.state_dict(),
            'scheduler': scheduler_gru.state_dict(),
            'train_loss': avg_train_loss,
            'val_loss': avg_val_loss,
            'best_val_loss': best_val_loss,
        }

```

```
    print(f"Saving new best model..")
    torch.save(state, 'model_gru.pth.tar')

    wandb.log({
        'train_loss': avg_train_loss,
        'val_loss': avg_val_loss,
        'lr': current_lr,
    })

wandb.finish()
print('Finished Training')
```

```
start training gru
```

```
Failed to detect the name of this notebook, you can set it manually with the  
WANDB_NOTEBOOK_NAME environment variable to enable code saving.  
wandb: Currently logged in as: demonstem. Use `wandb login --relogin` to force relogin
```

```
Tracking run with wandb version 0.16.4
```

```
Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern  
Recognition\HW5\wandb\run-20240320_162959-vym8rb1y
```

```
Syncing run curious-snowball-1 \(https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/vym8rb1y\) to Weights & Biases \(https://wandb.ai/demonstem/precipitation-nowcasting-gru\) (docs \(https://wandb.me/run\))
```

```
View project at https://wandb.ai/demonstem/precipitation-nowcasting-gru  
(https://wandb.ai/demonstem/precipitation-nowcasting-gru)
```

```
View run at https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/vym8rb1y  
(https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/vym8rb1y).
```

```
Training epoch 1...
```

```
Current LR: 0.01
```

```
Epoch 1 train loss: 1.9387
```

```
Validating epoch 1
```

```
Epoch 1 val loss: 1.6765
```

```
Saving new best model..
```

```
Training epoch 2...
```

```
Current LR: 0.01
```

```
Epoch 2 train loss: 1.9148
```

```
Validating epoch 2
```

```
Epoch 2 val loss: 1.6743
```

```
Saving new best model..
```

```
Training epoch 3...
```

```
Current LR: 0.01
```

```
Epoch 3 train loss: 1.9165
```

```
Validating epoch 3
```

```
Epoch 3 val loss: 1.6741
```

```
Saving new best model..
```

```
Training epoch 4...
```

```
Current LR: 0.01
```

```
Epoch 4 train loss: 1.9227
```

```
Validating epoch 4
```

Epoch 4 val loss: 1.6757

Training epoch 5...

Current LR: 0.01

Epoch 5 train loss: 1.9193

Validating epoch 5

Epoch 5 val loss: 1.6752

Training epoch 6...

Current LR: 0.01

Epoch 6 train loss: 1.9192

Validating epoch 6

Epoch 6 val loss: 1.6731

Saving new best model..

Training epoch 7...

Current LR: 0.01

Epoch 7 train loss: 1.9178

Validating epoch 7

Epoch 7 val loss: 1.6730

Saving new best model..

Training epoch 8...

Current LR: 0.01

Epoch 8 train loss: 1.9172

Validating epoch 8

Epoch 8 val loss: 1.6757

Training epoch 9...

Current LR: 0.01

Epoch 9 train loss: 1.9148

Validating epoch 9

Epoch 9 val loss: 1.6742

Training epoch 10...

Current LR: 0.002

Epoch 10 train loss: 1.9313

Validating epoch 10

Epoch 10 val loss: 1.6741

Run history:



Run summary:

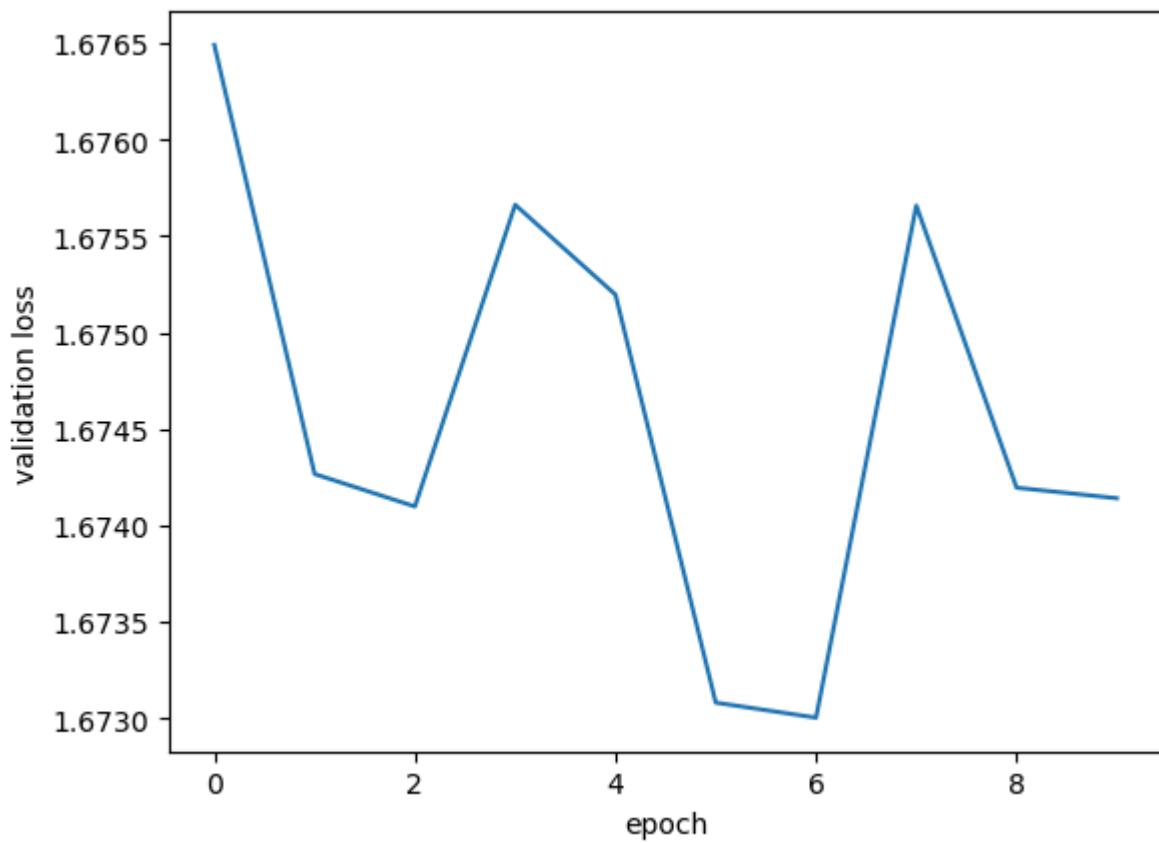
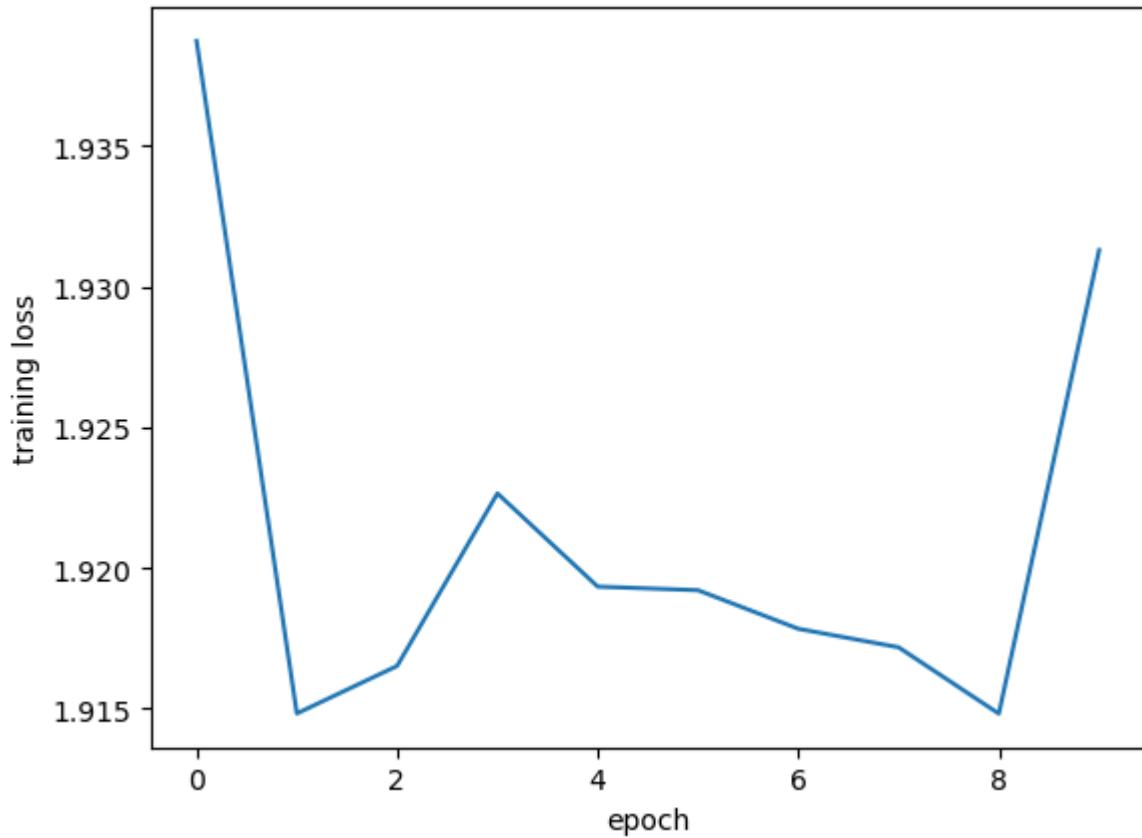
lr	0.002
train_loss	1.9313
val_loss	1.67414

View run **curious-snowball-1** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/vym8rb1y> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/vym8rb1y>)
Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_162959-vym8rb1y\logs

Finished Training

```
In [21]: # Plot
plt.plot(train_losses_gru)
plt.xlabel('epoch')
plt.ylabel('training loss')
plt.show()
plt.plot(val_losses_gru)
plt.xlabel('epoch')
plt.ylabel('validation loss')
plt.show()
```



```
In [16]: # Evaluate
def evaluate_gru(data_loader, model):
    """
    Evaluate model on validation data given by data_loader
    """

    # write code here
    # Start wandb run
    wandb.init(
        project='precipitation-nowcasting-gru',
    )

    # Log parameters and gradients
    wandb.watch(model, log='all')
    with torch.no_grad():
        train_losses = []
        for i, (inputs, y_true) in enumerate(tqdm(data_loader)):
            # Transfer data from cpu to gpu
            inputs = inputs.to(device)
            y_true = y_true.to(device)

            # Predict
            y_pred = model(inputs)

            # Calculate loss
            y_pred = torch.reshape(y_pred, (-1, 5))
            loss = loss_fn(y_pred, y_true)

            # Log stuff
            train_losses.append(loss)

        mse = torch.stack(train_losses).mean()
        wandb.finish()
    return mse
```

```
In [22]: checkpoint = torch.load('model_gru.pth.tar')
loaded_model = GRU(hidden_size=config_gru['hidden_size']) # Create model object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
model_gru = loaded_model
model_gru.to(device)
```

Loaded epoch 6 model

```
Out[22]: GRU(
    (gru): GRU(75, 200, num_layers=5)
    (ff1): Linear(in_features=200, out_features=200, bias=True)
    (out): Linear(in_features=200, out_features=1, bias=True)
)
```

```
In [25]: evaluated_train_gru = evaluate_gru(train_loader_gru, model_gru).item()
evaluated_val_gru = evaluate_gru(val_loader_gru, model_gru).item()

print('gru-model')
print('train', evaluated_train_gru)
print('validate', evaluated_val_gru)
```

Failed to detect the name of this notebook, you can set it manually with the WANDB_NOTEBOOK_NAME environment variable to enable code saving.

wandb: Currently logged in as: `demonstem`. Use `wandb login --relogin` to force relogin

Tracking run with wandb version 0.16.4

Run data is saved locally in `c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_170013-11822t8v`

Syncing run **comic-dream-3** (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/l1822t8v>) to [Weights & Biases](https://wandb.ai/demonstem/precipitation-nowcasting-gru) (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>) ([docs](https://wandb.me/run) (<https://wandb.me/run>))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-gru> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/l1822t8v> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/l1822t8v>)

View run **comic-dream-3** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/l1822t8v> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/l1822t8v>)
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: `.\wandb\run-20240320_170013-11822t8v\logs`

Tracking run with wandb version 0.16.4

Run data is saved locally in `c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_170140-cv0m1ltu`

Syncing run **twilight-field-4** (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cv0m1ltu>) to [Weights & Biases](https://wandb.ai/demonstem/precipitation-nowcasting-gru) (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>) ([docs](https://wandb.me/run) (<https://wandb.me/run>))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-gru> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cv0m1ltu> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cv0m1ltu>)

View run **twilight-field-4** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cv0m1ltu> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cv0m1ltu>)
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: `.\wandb\run-20240320_170140-cv0m1ltu\logs`

```
gru-model
train 1.9142119884490967
validate 1.673003077507019
```

Transformer

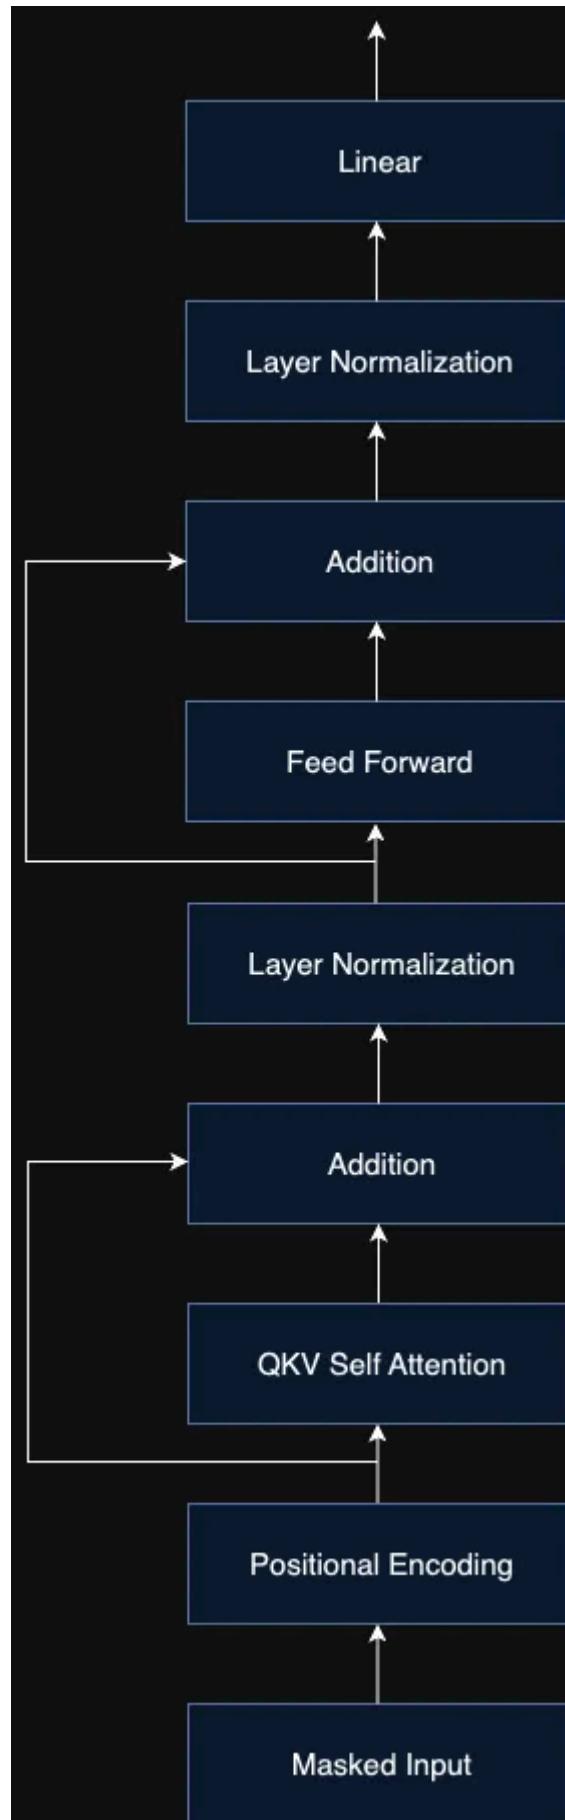
Welcome to the beginning of the real world! The above models are not usually used in practice due to its limited capability. Transformers are generally used by computer vision, natural language processing, and speech processing (almost every big AI fields).

In our dataloader, we will add the output of this timestep (the number of precipitation) as an auxiliary input to predict the next timestep. Thus, input for the model should be [#batch_size, 5, 76] (5 timesteps and the number 76 comes from $(3 \times 5 \times 5) + 1$) and the output for the model should be [#batch_size, 1] which would be the next timestep we want to predict. Additionally, we will mask the input at the dataloader to the attention from observing future values. Suppose that we want to predict timestep 3, we will mask the timestep 3, 4 and 5 in our input by setting it to zeros, and we will predict the timestep 3.

In order to get a score on this TODO, students need to implement a dataloader that mask the input correctly.

```
In [ ]: #####  
##  
# OT#2:  
#  
# Complete the code for preparing data for training Transformer  
#  
# Transformer's input should has 3 dimensions.  
#  
# The dimensions should compose of entries, time-step, and features.  
#  
#####  
##  
#                      WRITE YOUR CODE BELOW  
#  
#####  
##
```

In this task, we will implement one encoder layer of Transformer and add the linear layer to make a regression prediction. For the simplicity of the model, we will change the multi-head attention to QKV self-attention (single-head). As a result, our model should look like the diagram below. Since the layer self-attention is not available in torch, students have to implement it themselves. In Add & Norm layer, students have to do the addition before normalizing. In Layer Normalization, we will normalize across both timesteps and features.



```
=====
=====
Layer (type:depth-idx)           Output Shape      Param #
=====
=====
TransformerModel                  [1024, 1]          --
├PositionalEncoding: 1-1         [1024, 5, 76]     --
| └Dropout: 2-1                 [1024, 5, 76]     --
├SelfAttention: 1-2             [1024, 5, 76]     --
| └Linear: 2-2                  [1024, 5, 76]     5,852
| └Linear: 2-3                  [1024, 5, 76]     5,852
| └Linear: 2-4                  [1024, 5, 76]     5,852
| └Softmax: 2-5                 [1024, 5, 5]      --
├LayerNorm: 1-3                 [1024, 5, 76]     760
├Linear: 1-4                   [1024, 5, 76]     5,852
├LayerNorm: 1-5                 [1024, 5, 76]     760
└Linear: 1-6                   [1024, 1]          381
=====
=====
Total params: 25,309
Trainable params: 25,309
Non-trainable params: 0
Total mult-adds (M): 25.92
=====
=====
Input size (MB): 1.56
Forward/backward pass size (MB): 18.69
Params size (MB): 0.10
Estimated Total Size (MB): 20.34
=====
```

```
In [ ]: #####  
##  
# OT#3  
#  
# Write a PyTorch PositionalEncoding model.  
#  
#  
#  
# Hint: You should read PyTorch documentation to see the list of available  
#  
# layers and options you can use.  
#  
#####  
##  
#  
#  
# WRITE YOUR CODE BELOW  
#  
#####  
##  
  
class PositionalEncoding(nn.Module):  
    def __init__(self, seq_len, emb_dim, dropout=0.2):  
        pass  
  
    def forward(self, x):  
        pass
```

```
In [ ]: #####  
##  
# OT#4  
#  
# Write a PyTorch Transformer model.  
#  
# Your goal is to predict a precipitation of every time step.  
#  
#  
# Hint: You should read PyTorch documentation to see the list of available  
#  
# Layers and options you can use.  
#  
#####  
##  
# WRITE YOUR CODE BELOW  
#  
#####  
##
```



```
class SelfAttention(nn.Module):  
    def __init__(self, input_dim):  
        pass  
  
    def forward(self, x):  
        pass  
  
class TransformerModel(nn.Module):  
    def __init__(self):  
        pass  
  
    def forward(self, x):  
        pass
```

```
In [ ]: #####  
##  
# OT#5  
#  
# Complete the code to train your Transformer model  
#  
#####  
##  
print('start training transformer')  
#####  
##  
# WRITE YOUR CODE BELOW  
#  
#####  
##
```

```
In [ ]: # Plot
```

If you implement it correctly, you should evaluate the model in the test dataset and the score should be better than the above models.

```
In [ ]: # Evaluate
```

Final Section

PyTorch playground

Now, train the best model you can do for this task. You can use any model structure and function available. Remember that training time increases with the complexity of the model. You might find printing computation graphs helpful in debugging complicated models.

Your model should be better than your CNN or GRU model in the previous sections.

Some ideas:

- Tune the hyperparameters
- Adding dropouts
- Combining CNN with GRUs

You should tune your model on training and validation set.

The test set should be used only for the last evaluation.

```
In [ ]: # Prep data as you see fit
```

In [109]:

```
#####
## 
# TODO#20
#
# Write a function that returns your best PyTorch model. You can use anything
#
# you want. The goal here is to create the best model you can think of.
#
#
#
# Hint: You should read PyTorch documentation to see the list of available
#
# Layers and options you can use.
#
#####
## 
#                               WRITE YOUR CODE BELOW
#
#####
## 
```

```
class Fancy(nn.Module):
    def __init__(self, hidden_size=200):
        super(Fancy, self).__init__()
        self.gru = nn.GRU(3 * 5 * 5, hidden_size, 5)
        self.ff1 = nn.Linear(hidden_size, hidden_size)
        self.ff2 = nn.Linear(hidden_size, hidden_size)
        self.ff3 = nn.Linear(hidden_size, hidden_size)
        self.dropout = nn.Dropout(0.5)
        self.out = nn.Linear(hidden_size, 1)

    def forward(self, x):
        x = self.gru(x)
        x = F.relu(self.ff1(x[0]))
        x = self.dropout(x)
        x = F.relu(self.ff2(x))
        x = self.dropout(x)
        x = F.relu(self.ff3(x))
        y = self.out(x)
        return y
```

```
In [110]: fancy_config = {
    'architecture': 'gru',
    'lr': 0.01,
    'hidden_size': 100,
    'scheduler_factor': 0.2,
    'scheduler_patience': 2,
    'scheduler_min_lr': 1e-4,
    'epochs': 5
}

# Model
model_fancy = Fancy(hidden_size=fancy_config['hidden_size'])
model_fancy = model_fancy.to(device)
optimizer_fancy = torch.optim.Adam(model_fancy.parameters(), lr=fancy_config['lr'])
scheduler_fancy = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer_fancy,
    'min',
    factor=fancy_config['scheduler_factor'],
    patience=fancy_config['scheduler_patience'],
    min_lr=fancy_config['scheduler_min_lr']
)
```

```
In [ ]: #####  
##  
# TODO#21  
#  
# Complete the code to train your best model  
#  
#####  
##  
print('start training the best model')  
#####  
##  
# WRITE YOUR CODE BELOW  
#  
#####  
##  
  
# Training  
train_losses_fancy = []  
val_losses_fancy = []  
learning_rates_fancy = []  
  
# Start wandb run  
wandb.init(  
    project='precipitation-nowcasting-dropout',  
    config=fancy_config,  
)  
  
# Log parameters and gradients  
wandb.watch(model_fancy, log='all')  
  
for epoch in range(fancy_config['epochs']): # Loop over the dataset multiple times  
  
    # Training  
    train_loss = []  
    current_lr = optimizer_fancy.param_groups[0]['lr']  
    learning_rates_fancy.append(current_lr)  
  
    # Flag model as training. Some layers behave differently in training and  
    # inference modes, such as dropout, BN, etc.  
    model_fancy.train()  
  
    print(f"Training epoch {epoch+1}...")  
    print(f"Current LR: {current_lr}")  
  
    for i, (inputs, y_true) in enumerate(tqdm(train_loader_gru)):  
        # Transfer data from cpu to gpu  
        inputs = inputs.to(device)  
        y_true = y_true.to(device)  
  
        # Reset the gradient  
        optimizer_fancy.zero_grad()  
  
        # Predict  
        y_pred = model_fancy(inputs)
```

```

# Calculate loss
y_pred = torch.reshape(y_pred, (-1, 5))
loss = loss_fn(y_pred, y_true)

# Compute gradient
loss.backward()

# Update parameters
optimizer_fancy.step()

# Log stuff
train_loss.append(loss)

avg_train_loss = torch.stack(train_loss).mean().item()
train_losses_fancy.append(avg_train_loss)

print(f"Epoch {epoch+1} train loss: {avg_train_loss:.4f}")

# Validation
model_fancy.eval()
with torch.no_grad(): # No gradient is required during validation
    print(f"Validating epoch {epoch+1}")
    val_loss = []
    for i, (inputs, y_true) in enumerate(tqdm(val_loader_gru)):
        # Transfer data from cpu to gpu
        inputs = inputs.to(device)
        y_true = y_true.to(device)

        # Predict
        y_pred = model_fancy(inputs)

        # Calculate loss
        y_pred = torch.reshape(y_pred, (-1, 5))
        loss = loss_fn(y_pred, y_true)

        # Log stuff
        val_loss.append(loss)

    avg_val_loss = torch.stack(val_loss).mean().item()
    val_losses_fancy.append(avg_val_loss)
    print(f"Epoch {epoch+1} val loss: {avg_val_loss:.4f}")

# LR adjustment with scheduler
scheduler_fancy.step(epoch)

# Save checkpoint if val_loss is the best we got
best_val_loss = np.inf if epoch == 0 else min(val_losses_fancy[:-1])
if avg_val_loss < best_val_loss:
    # Save whatever you want
    state = {
        'epoch': epoch,
        'model': model_fancy.state_dict(),
        'optimizer': optimizer_fancy.state_dict(),
        'scheduler': scheduler_fancy.state_dict(),
        'train_loss': avg_train_loss,
        'val_loss': avg_val_loss,
    }

```

```
        'best_val_loss': best_val_loss,
    }

    print(f"Saving new best model..")
    torch.save(state, 'model_fancy4.pth.tar')

wandb.log({
    'train_loss': avg_train_loss,
    'val_loss': avg_val_loss,
    'lr': current_lr,
})

wandb.finish()
print('Finished Training')
```

In [111]:

```
checkpoint = torch.load('model_fancy2.pth.tar')
loaded_model = Fancy(hidden_size=fancy_config['hidden_size']) # Create model object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
model_fancy = loaded_model
model_fancy.to(device)
```

Loaded epoch 0 model

Out[111]:

```
Fancy(
    (gru): GRU(75, 100, num_layers=5)
    (ff1): Linear(in_features=100, out_features=100, bias=True)
    (ff2): Linear(in_features=100, out_features=100, bias=True)
    (ff3): Linear(in_features=100, out_features=100, bias=True)
    (dropout): Dropout(p=0.5, inplace=False)
    (out): Linear(in_features=100, out_features=1, bias=True)
)
```

```
In [112]: # Evaluate best model on validation and test set
evaluated_test_fancy = evaluate_gru(test_loader_gru, model_fancy).item()
print('my-best-model', evaluated_test_fancy)
```

Finishing last run (ID:cl5ibdo8) before initializing another...

View run **silvery-leaf-12** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cl5ibdo8> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/cl5ibdo8>)

Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_232633-cl5ibdo8\logs

Successfully finished last run (ID:cl5ibdo8). Initializing new run:

Tracking run with wandb version 0.16.4

Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_232934-sbzd3n4g

Syncing run **atomic-night-13** (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/sbzd3n4g>) to Weights & Biases (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>) ([docs](#) (<https://wandb.me/run>))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-gru> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>).

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/sbzd3n4g> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/sbzd3n4g>)

View run **atomic-night-13** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/sbzd3n4g> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/sbzd3n4g>)

Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_232934-sbzd3n4g\logs

my-best-model 1.1568195819854736

```
In [68]: # Also evaluate your fully-connected model and CNN/GRU/Transformer model on the test set.
checkpoint = torch.load('model_cnn.pth.tar')
loaded_model = CNN(hidden_size=config_cnn['hidden_size']) # Create model object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
model_cnn = loaded_model
model_cnn.to(device)

checkpoint = torch.load('model_gru.pth.tar')
loaded_model = GRU(hidden_size=config['hidden_size']) # Create model object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
model_gru = loaded_model
model_gru.to(device)

checkpoint = torch.load('model_ff.pth.tar')
loaded_model = FeedForwardNN(hidden_size=config_gru['hidden_size']) # Create model object
loaded_model.load_state_dict(checkpoint['model']) # Load weights
print(f"Loaded epoch {checkpoint['epoch']} model")
model_ff = loaded_model
model_ff.to(device)
```

```
Loaded epoch 2 model
Loaded epoch 6 model
Loaded epoch 4 model
```

```
Out[68]: FeedForwardNN(
    (ff1): Linear(in_features=75, out_features=200, bias=True)
    (ff2): Linear(in_features=200, out_features=200, bias=True)
    (ff3): Linear(in_features=200, out_features=200, bias=True)
    (out): Linear(in_features=200, out_features=1, bias=True)
)
```

```
In [69]: evaluated_test_cnn = evaluate_cnn(test_loader, model_cnn).item()
print('cnn-model', evaluated_test_cnn)
```

Finishing last run (ID:wxzbxqxb) before initializing another...

View run **glad-aardvark-9** at: <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wxzbzbqxb> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wxzbzbqxb>)
Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_213133-wxzbzbqxb\logs

Successfully finished last run (ID:wxzbxqxb). Initializing new run:

Tracking run with wandb version 0.16.4

Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_213329-wmx4g350

Syncing run **major-armadillo-10** (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wmx4g350>) to [Weights & Biases](https://wandb.ai/demonstem/precipitation-nowcasting-cnn) (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>) ([docs](https://wandb.me/run) (<https://wandb.me/run>))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wmx4g350> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wmx4g350>)

View run **major-armadillo-10** at: <https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wmx4g350> (<https://wandb.ai/demonstem/precipitation-nowcasting-cnn/runs/wmx4g350>)
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_213329-wmx4g350\logs

cnn-model 1.161057710647583

```
In [72]: evaluated_test_ff = evaluate(test_loader, model_ff).item()
print('ff-model', evaluated_test_ff)
```

ff-model 1.1636040210723877

```
In [74]: evaluated_test_gru = evaluate_gru(test_loader_gru, model_gru).item()
print('gru-model', evaluated_test_gru)
```

Finishing last run (ID:epualq7v) before initializing another...

View run **exalted-surf-5** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/epualq7v> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/epualq7v>)
Synced 6 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_213851-epualq7v\logs

Successfully finished last run (ID:epualq7v). Initializing new run:

Tracking run with wandb version 0.16.4

Run data is saved locally in c:\Users\Tonza\Desktop\Code\Pattern Recognition\HW5\wandb\run-20240320_214042-1iooyesd

Syncing run **different-plasma-6** (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/1iooyesd>) to [Weights & Biases](#) (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>) ([docs](#) (<https://wandb.me/run>))

View project at <https://wandb.ai/demonstem/precipitation-nowcasting-gru> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru>)

View run at <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/1iooyesd> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/1iooyesd>)

View run **different-plasma-6** at: <https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/1iooyesd> (<https://wandb.ai/demonstem/precipitation-nowcasting-gru/runs/1iooyesd>)
Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)

Find logs at: .\wandb\run-20240320_214042-1iooyesd\logs

gru-model 1.1548452377319336

To get full credit for this part, your best model should be better than the previous models on the **test set**.

TODO#22

Explain what helped and what did not help here

Ans: add more hidden layer does not help much, what I did are adding dropouts, add more layers and decrease lr, and epoch for gru model. But got 1.156 > 1.154(old gru), therefore adding dropouts in this model does not help much either.

[Optional] Augmentation using data loader

Optional TODO#6

Implement a new dataloader on your best model that will perform data augmentation. Try adding noise of zero mean and variance of $10e^{-2}$.

Then, train your model.

```
In [ ]: # Write Dataset/DataLoader with noise here
```

```
In [ ]: print('start training the best model with noise')
#####
##                                     WRITE YOUR CODE BELOW
#
#####
##
```

```
In [ ]: # Evaluate the best model trained with noise on validation and test set
```